



- Home
- Wireshark cheat sheet
- Kail Pentest Netsec
- Rust
- ▼ Ansible
- Crystal
- ▼ Docker
- Elastic
- KDB Cheatsheet
- ▼ Git
- ▼ Golang
- JavaScript
- ▲ Linux
  - Iptables Basics
  - Linux Cheat Sheet**
  - OpenVPN - Client config
  - Hi Frequency/Volume Trading - OS Tuning
  - SSH Certificate-based Authentication
  - Troubleshooting frozen system
  - RAID levels
  - Barrier - screen control across physical devices
  - QEMU Virtualization
  - TCP troubleshooting
  - Marten web framework
  - Graylog
  - htmx & hyperscript
  - UDP packet loss troubleshooting
  - SSD types and terminology
  - Data Structures
- ▼ More...
- Julia cheatsheet
- Sonicwall
- Ninja Tools
- ▼ Python
- ▼ Ruby
- Ruby Cheat Sheet
- ▼ SaltStack
- Scratchpad
- ▼ Vagrant
- Vue JS Cheatsheet

# Linux Cheat Sheet

## BASH Internals

```
syntax, use 'echo' + keyword

$$ - PID of current shell
$0 - show shell name
$! - PID of last background cmd
$? - exist status of last cmd
$_ - previously created dir (mkdir foo && cd $_)
#@ - show all command's parameters
 $# - show # of arguments passed to command
$* - All arguments passed to command
$1 - first argument passed to command
!! - run previous command
-eq - math equal (int)
-ne - math not equal (int)
-lt - math less than
-le - math less than or equal
-gt - math greater than
-ge - math greater than or equal
-z - string is 0 length (null)

-n - string is not 0 length (not null)

        if [ -n "{var}" ] # if not null then..

-nt - newer than (file or object time)
(-r,-w,-x) - if object is readable,writable,exec
```

```
set default value if parameter is null or empty

var1=""
echo "${var1:-abc}"
>> abc
```

```
error out if parameter is null or empty

var1=""
echo "${var1:?}"
>>  bash: var1: parameter null or not set
```

```
check if character is inside a string

str="1,2,3,4-5"
[[ "$str" == *-* ]] && echo "has dash"
```

```
source file in same directory as caller script
source "$(readlink -f $0 | xargs dirname)/shared.txt"
```

## Arrays & Dictionaries

```
Simple Array

Fruits=('Apple' 'Banana' 'Cherry')
echo ${Fruits[1]} ## Banana
echo ${Fruits[@]:1:2} ## range from 2nd to 3rd element,  banana cherry
```

```
add an element to array

Fruits+=('Watermelon')
echo ${Fruits[@]} ## "apple, banana, cherry, watermelon",  @ = show all elements in array
```

```
Loop over an array

arr=(apples oranges tomatoes)
# Just elements.
for element in "${arr[@]}; do
    printf '%s\n' "$element"
done
```

```
create array from string with delimiter

str="1,2,3,4,5"
IFS="," read -a ARR <<< $str
echo "${ARR[@]}"
>>> 1 2 3 4 5
```

## Key/Value pairs (associative Array, aka Hash, Dictionary)

```
read in config file, check array of key/val pairs to make sure parameters are set

$config = "/etc/file.conf"
declare -A mylist=( [first]=
                    [last]=
                    [age]=
                    )

for param in "${!myList[@]}"; do
    value=$(grep ^$param $config)
    var[$param]=$value

    if [ -z ${myList[$param]} ]; then
        echo "param $param is not set"; exit 1
    fi
done
```

## Dictionary/Hash in bash (parse IP + Port hash, netcat to each IP and port)

```
readonly connections=(

'A, 205.209.202.37, 7755'
'B, 205.209.202.1, 8899'
'C, 205.209.202.21, 4578'

)

function nctest(){
    local name ip port
    for fields in ${connections[@]}
    do

        # strip whitespace
        fields=$(echo -e "${fields}" | tr -d '[:space:]')

        IFS=$', ' read -r typ name ip port <<< $fields

        conn=$(nc -zv -w 2 $ip $port 2>&1 | grep 'Connection refused' )

        if [[ -z "$conn" ]] then
            echo "[name] nc $ip $port OK"
        else
            echo "[name] nc $ip $port REFUSED"
        fi

        echo "-----"
        done
    }

nctest
```

```
basic dict

declare -A sounds
sounds[dog]="bark"
sounds[cow]="moo"
sounds[cat]="meow"
echo ${sounds[dog]} # bark
```

```
Iterate keys and values

for val in "${sounds[@]}"; do
    echo $val
done
```

```
for key in "${!sounds[@]}"; do
    echo $key
done
```

## Math

```
calculate value

echo $((35+15))
50
```

```
generate random number 0 to 500

$((RANDOM % 500))
```

```
use calculator

echo "12+3" | bc # 15
echo "10^2" | bc # 100
echo "10/2" | bc # 5
```

```
add 2 variables

sum=$((var1+var2))
```

## Loops and Conditionals

```
for loop with Range

for loop in {1..50}; do echo "processing $loop"; sleep 2; done
> processing 1
> processing 2
etc
```

```
Range with step size, count every 5

for i in {1..50..5}
```

```
or use a 'seq' operand

for i in $(seq 1 5); do echo $i; done
1
2
3..etc
```

```
sequence with step size,

for i in $(seq 1 5 30); do echo $i; done
1
6
11
16
21
26
```

```
Counter loop

for ((i=0; i<100; i++)); do echo $i; done
```

```
loop thru directories and grep something from config files

for i in $(ls); do grep 'something' $i/*.conf ; done
```

```
Case statement (check input params)

case $key in
-u| -username | --username)
    UNAME="$2"
    shift
    ;;
-pw| -password | --password)
    PASSWORD="$2"
    shift
    ;;
-p| -profile | --profile)
    PROFILE="$2"
    shift
    ;;
*)
    echo "Unknown Option"
    exit 1
```

```
case $env in
    "test")
        username="testuser"; password="testPW";;
    "prod")
        username="produser"; password="prodPW";;
esac
```

```
If statement

if [[ -z "$string" ]];
then
    echo "String is empty"
elif [[ -n "$string" ]]; then
    echo "String is not empty"
fi
```

```
Variable Conditionals

[[ -z STRING ]] Empty string
[[ -n STRING ]] Not empty string
[[ STRING == STRING ]] Equal
[[ STRING != STRING ]] Not Equal
[[ NUM -eq NUM ]] Equal
[[ NUM -ne NUM ]] Not equal
[[ NUM -lt NUM ]] Less than
[[ NUM -le NUM ]] Less than or equal
[[ NUM -gt NUM ]] Greater than
[[ NUM -ge NUM ]] Greater than or equal
[[ STRING =~ STRING ]] Regexp
(( NUM < NUM )) Numeric conditions
[[ -o noclobber ]] If OPTIONNAME is enabled
[[ ! EXPR ]] ## "Not something"
[[ X ]] && [[ Y ]] ## X And Y
[[ X ]] || [[ Y ]] ## X Or Y
```

```
File conditionals

[[ -e FILE ]] Exists
[[ -r FILE ]] Readable
[[ -h FILE ]] Symlink
[[ -d FILE ]] Directory
[[ -w FILE ]] Writable
[[ -s FILE ]] Size is > 0 bytes
[[ -f FILE ]] File is type "file"
[[ -x FILE ]] Executable
[[ FILE1 -nt FILE2 ]] 1 is more recent than 2
[[ FILE1 -ot FILE2 ]] 2 is more recent than 1
[[ FILE1 -ef FILE2 ]] Same files
```

```
Case / Switch

case "$1" in
    start | begin)
        service start
    ;;
    stop | kill)
        service stop
    ;;
    *)
        echo "usage: $0 {start|stop}"
    ;;
esac
```

## GREP

```
search for joe in names.txt
grep 'joe' /names.txt
```

```
search for 'joe' in directory dir1
grep 'joe' /dir1 -r
```

```
search for 'joe' in dir1 and follow symlinks
grep 'joe' /dir1 -R
```

```
search only files that match
grep 'joe' /dir1 -l
```

```
search only files that dont match
grep 'joe' /dir1 -L
```

```
case insensitive
grep 'JOE' /dir1 -i
```

```
add col
```

This site uses cookies from Google to deliver its services and to analyze traffic. Information about your use of this site is shared with Google. By using this site, you agree to its use of cookies.



mrxpalmeiras

- Home
- Wireshark cheat sheet
- Kail Pentest Netsec
- Rust
- Ansible
  - Crystal
- Docker
  - Elastic
  - KDB Cheatsheet
- Git
- Golang
- JavaScript
- Linux
  - Iptables Basics
  - Linux Cheat Sheet
  - OpenVPN - Client config
  - Hi Frequency/Volume Trading - OS Tuning
  - SSH Certificate-based Authentication
  - Troubleshooting frozen system
  - RAID levels
  - Barrier - screen control across physical devices
  - QEMU Virtualization
  - TCP troubleshooting
  - Marten web framework
  - Graylog
  - htmx & hyperscript
  - UDP packet loss troubleshooting
- SSD types and terminology
- Data Structures
- More...
  - Julia cheatsheet
  - Sonicwall
  - Ninja Tools
  - Python
  - Ruby
  - Ruby Cheat Sheet
  - SaltStack
  - Scratchpad
  - Vagrant
  - Vue JS Cheatsheet

## Debugging & Test

enable line by line processing output in script

```
set -x
```

verbose output only if debug flag is set

```
debug=1
```

```
test $debug -gt 0 && echo "var is $var"
```

test if file exists, if not, exit with error

```
test -f "${config}" || { echo "${config} not present, exiting.."; exit 1; }
```

check if parameters are set and not empty, exits out w error if not set

```
err_msg="[ERROR] parameter is not set or empty value:"
myParam=${1:?}$err_msg You have an error, missing myParam"
```

check segfault core output

```
gdb $(which python) core.python.$date
gdb <my binary> <path to core>
```

## Regex & String Manipulation

convert uppercase files to lowercase

```
rename 'y/A-Z/a-z/' *
```

capitalize 1st letter

```
var="wunderbar"
echo ${var^^}
```

capitalize entire word

```
echo ${var^^}
```

get # of characters in variable

```
var="milkshake"
echo ${#var}
9
```

Check if Word is in a String

```
[[ "$string" == *"$word"* ]] || echo "word not in string"
```

strip off the last character from a string,

```
var="Banana"
echo ${var%?} // Banana
```

Get value between 2 delimiters,

```
grep ExecStart bitbucket.service | awk -v FS="(bitbucket/|bin)" '{print $2}'
```

extract filename from a path

```
echo /somedir/blah/postgresql96-9.6.5.x86_64.rpm | awk '{match($1, "[^/]*$", a)}END{print a[0]}'
postgresql96-9.6.5.x86_64.rpm
```

search for a pattern in all files

```
grep -RnisI "My cat*" /var/log/*
```

Insert string after a delimiter, save in place (insert "dog" after "cat")

```
tmpfile=$(mktemp)
awk '/cat/ { print; print "dog"; next}1' pets.txt > $tmpfile && mv -f $tmpfile pets.txt
```

String manipulation

```
STR="HELLO WORLD!"
echo ${STR,,} #=> "hello world!" (lowercase 1st letter)
echo ${STR,,,} #=> "hello world!" (all lowercase)
STR="hello world!"
echo ${STR^} #=> "Hello world!" (uppercase 1st letter)
echo ${STR^^} #=> "HELLO WORLD!" (all uppercase)

# Substitution

VAR="beachball"

${VAR%suffix} # Remove suffix
${VAR%prefix} # Remove prefix
${VAR%suffix} # Remove long suffix
${VAR##prefix} # Remove long prefix
${VAR/beach/basket} # Replace first match >> basketball

## Length of string
echo ${#VAR} # 9
```

grep for multiple strings

```
ls -la /home | grep -v "joe\|fred\|bob"
```

another way,

```
cat /etc/passwd | egrep -v '^(root|halt|sync|shutdown|adm|bin|daemon)'
```

get last character of a string

```
permissions="775"
echo "${permissions: -1}"
```

get 2nd and 3rd character from a string

```
str="boris"
second=$(echo $str | head -c 2 | tail -c 1)
third=$(echo $str | head -c 3 | tail -c 1)
```

check if word is in a string

```
str="sun is shining"
[ -z "${str##*'shining'*}" ] && echo "contains word!!!"
```

break down string using delimiter

```
str="deny=5"
key=${str%%=*} # key is 'deny'
val=${str##*=} # val is 5
```

### Cut command

breakdown string by fields

```
string="/mnt/hc/home/user"
echo $string | cut -d '/' -f2 ## mnt
echo $string | cut -d '/' -f2,3,4 ## mnt hc home

# from 2nd field to end of string
echo $string | cut -d '/' -f2- ## mnt/hc/home/user
```

## Shell Cmds

run python inside Bash with arguments

```
function print_hello {
    NAME="${1}" python - <<END
    import os
    print("Hello there %s" % os.environ['NAME'])
    END
}

print_hello Joe
```

get variable from a json dump using python

```
URL=$(echo ${URL} | python -c 'import sys,json; print json.load(sys.stdin)["url"]')
```

colorize Bash prompt (insert into ~/.bashrc)

```
export PS1="\[\e[31m\]\u\[\e[m\]\[\e[33m\]@\h\[\e[m\]:\W\]$ "
```

Root PS1

```
export PS1="\[\e[30;41m\]\u\[\e[m\]\[\e[33m\]@\h\[\e[m\]:\W\]$ "
```

generate a random password

```
date +%s | sha256sum | base64 | head -c 8; echo
```

remove all empty directories

```
find . -type d -empty -delete
```

copy permissions on file1 to file2

```
chmod --reference file2 file1
```

remove all but specific file

```
rm -f !(theFile.txt)
```

remove files that dont match a specific extension

```
rm !(*.xls|*.xlsx|*.csv)
```

find duplicate files (check file hash)

```
find -not -empty -type f -printf "%s
" | sort -rn | uniq -d | xargs -I{} -n1 find -type f -size {}c -print0 |
xargs -0 md5sum | sort | uniq -w32 --all-repeated=separate
```

run a command as another user

```
runuser -l joe -c 'whoami'
```

## JSON & YAML

### JQ - json parser

show all values in PP format

```
jq . file.json
```

PP api output

```
curl example.org/api/v1/users | jq .
```

show specific key

```
json={"name":"","bob", "age":23}
echo $json | jq '.name'
```

parse array key for specific value

```
echo $json | jq '.values[].title'
```

select multiple properties of 1st element

```
jq '.[0] | { _id, email }' file.json
```

delete key

```
cat file.json
{ "name": "joe", "age": 23, "user-name": "j123" }

delete Name key
jq 'del(.name)' file.json > file2.json

delete key with dashes
jq 'del(.user-name)' file.json > file2.json
```

Convert YAML to JSON - 1 liner

```
python -c 'import sys, yaml, json; json.dump(yaml.load(sys.stdin), sys.stdout, indent=4)' < file.yaml > file.json
```

### Functions

basic function

```
myFunc() {
    echo "hello $1"
}

myFunc "bob"
```

### Kill

```
kill -1 PID # SIGHUP, shutdown proc + restart
kill -2 PID # TERM, same as control+c
kill -3 PID # CORE, stop proc, create a core dump
kill -9 PID # SIGKILL, kill unresponsive proc, dirty kill
kill -11 PID # SIGSEGV, create core dump on segmentation fault, useful for misbehaving procs
kill -15 PID # TERM, default kill flag, same as "kill PID"
```

This site uses cookies from Google to deliver its services and to analyze traffic. Information about your use of this site is shared with Google. By using this site, you agree to its use of cookies.

LEARN MORE GOT IT



mrxpalmeiras

- Home
- Wireshark cheat sheet
- Kail Pentest Netsec
- Rust
- Ansible
- Crystal
- Docker
- Elastic
- KDB Cheatsheet
- Git
- Golang
- JavaScript
- Linux
  - Iptables Basics
  - Linux Cheat Sheet
  - OpenVPN - Client config
  - Hi Frequency/Volume Trading - OS Tuning
  - SSH Certificate-based Authentication
  - Troubleshooting frozen system
  - RAID levels
  - Barrier - screen control across physical devices
  - QEMU Virtualization
  - TCP troubleshooting
  - Marten web framework
  - Graylog
  - htmx & hyperscript
  - UDP packet loss troubleshooting
- SSD types and terminology
- Data Structures
- More...
  - Julia cheatsheet
  - Sonicwall
  - Ninja Tools
  - Python
  - Ruby
  - Ruby Cheat Sheet
  - SaltStack
  - Scratchpad
  - Vagrant
  - Vue JS Cheatsheet

## Various

add commands alias  
alias ls='ls -lta --color=auto'

Run a specific cmd from history

```
history
120 cat /var/log/messages
121 vi /etc/hosts
!120
## will show /var/log/messages
```

redirect std output to both file and screen  
program [arguments...] 2>&1 | tee outfile

Get Date in specific format  
echo \$(date +%Y%m%d\_%H%M%S)

```
Parameter Expansion

name="John"
echo ${name}
echo ${name//j/} #=> "john" (substitution)
echo ${name:0:2} #=> "Jo" (slicing)
echo ${name::2} #=> "Jo" (slicing)
echo ${name::-1} #=> "Joh" (slicing)
echo ${name:(-1)} #=> "n" (slicing from right)
echo ${name:(-2):1} #=> "h" (slicing from right)
echo ${food:-Cake} #=> $food or "Cake"

length=2
echo ${name:0:length} ## Jo

file="/mnt/dir/abc.cpp"
echo ${basename $file} # abc.cpp (basepath)
echo ${file##*/} # foo.cpp (basepath)
echo ${file%.*} ## /mnt/dir/abc
echo ${file###*.} # cpp (extension)

echo ${STR#*/} # path/to/foo.cpp
echo ${STR##*/} # foo.cpp
echo ${STR/foo/bar} # /path/to/bar.cpp

## set default values

${F00:-val} ## $F00, or val if not set
example: port=${1:-22} # if port not set via argument, make it 22

${F00:=val} ## Set $F00 to val if not set
${F00:+val} ## val if $F00 is set
${F00:?message} ## Show error message and exit if $F00 is not set
```

## SSH

```
file permissions

/home/user = 700
/home/user/.ssh = 700
/home/user/.ssh/id_rsa = 600
/home/user/.ssh/id_rsa.pub = 644
/home/user/.ssh/authorized_keys = 600
/home/user/.ssh/known_hosts = 644
```

troubleshoot auth errors  
on target (where youre trying to ssh into), start SSH on different port, debug mode  
/usr/sbin/sshd -d -p 2222

on client, connect to target  
ssh user@target -p 2222 -vvv

SSH Shuttle  
pip3 install sshuttle

route all connections to 172.31.23.156 via "server B"  
sshuttle -r user@<server B IP> 172.31.23.156

all connections will now be going via remote IP, encrypted  
to route ALL connections, use 0/0  
sshuttle -r user@serverB 0/0

proxy connections for a specific website,via jump host serverB, send to background  
serverA> nohup sshuttle -r serverB `dig +short www.somesite.com | sed "/[^0-9\\.]/d"  
| xargs -n1 -I '\$' echo -n '\$/32 '` 2>&1 &

pass a custom SSH key to sshuttle  
sshuttle --dns user@host <IP range> --ssh-cmd 'ssh -i /home/user/priv\_key'

use SSH as a web proxy  
ssh -D 8080 username@proxyHost

set browsers proxy option to 127.0.0.1:8080, all browsing requests will go via proxyHost

TMUX  
start tmux session  
tmux

reattach to session after broken connection  
tmux ls  
0: 1 windows (created Tue Aug 23 12:39:52 2011) [103x30]  
tmux attach -t 0

delete session  
tmux kill-session -t 0

File Operations  
create 25 new files from one command, use: {1..X}  
touch myfile{1..25}

get file extension  
file=superman.jpg  
name=\${file%.\*} # superman  
ext=\${file#\*.} # jpg

delete all files that dont match an extension  
rm !(\*.foo|\*.bar|\*.baz)

delete files from search  
find . -name '\*.pyc' -delete

grep 5 lines above and below a certain value  
cat employees.txt | grep -A 5 -B 5 'Mr. Jones'

remove all blank lines from a file  
grep . file1 > file2

read in a file  
< file.txt | while read line; do  
echo \$line  
done

generate 1GB empty file  
dd if=/dev/zero of=testfile count=1024000 bs=1024

or  
fallocate -l 1GB testfile

create a random large 200MB file,  
dd if=/dev/urandom of=file.txt bs=2075200 count=100

generate 10mb file with random text  
base64 /dev/urandom | head -c 10000000 > testfile

VIM  
delete all lines from file  
:1,\$d

search for all instances of string 'horse'  
escape key  
/horse  
press 'n' to move to next occurrence

vi a file on remote server  
vi scp://user@<hostname>//etc/hosts

check what pub key matches the priv key  
ssh-keygen -y -e -f ~/.ssh/id\_rsa

add a new SSH key and copy the public key to remote known\_hosts file  
ssh-keygen -t rsa  
cat ~/.ssh/id\_rsa.pub | ssh user@hostname 'cat >> .ssh/authorized\_keys'

run a command on remote host  
ssh servername cmdname

connect to an unreachable server B (port 2345) via SSH hop over reachable server A  
ssh user@serverA -L 6789:serverB:2345 -f -N (localhost:6789 = serverB:2345)

Port Tunneling via SSH  
(port 1200 is unreachable from server A, connect to it via localhost:1300 via SSH to server B  
user@serverA> ssh -L 1300:localhost:1200 serverB -fN

or via SSH Jumping  
A > B > C (B has to have AllowTCPForwarding=yes in sshd\_config)  
A> ssh -J user@B user@C

Proxy to an unreachable server via reachable  
A can talk to B  
B can talk to C  
A cant talk to C (but needs to)  
user@serverB> ssh -L 0.0.0.0:9222:serverC:22 (will SSH into C)  
user@serverA> ssh serverB -p 9222 (will ssh you into C)

```
setup SSH Sockets
mkdir ~/.ssh/sockets

vim ~/.ssh/config
UseRoaming no
TCPKeepAlive yes
ServerAliveInterval 15
ServerAliveCountMax 6
Host *
Compression yes
ControlMaster auto
ControlPath ~/.ssh/sockets/%r@h:%p
ControlPersist yes
ControlPersist 600

Host nycweb1
Hostname 192.168.10.2
User root
IdentityFile ~/.ssh/id_rsa
```

show fingerprint of a public key file, useful to track down /var/log/secure messages to see who logged in  
ssh-keygen -lf /home/user/.ssh/authorized\_keys | grep <fingerprint> (looks like SHA256:zZUd2W)

Use a hop server to access a unreachable host, add to ~/.ssh/config  
Host <target-host>  
ProxyCommand ssh -q -W %h:%p <hop-host>

This site uses cookies from Google to deliver its services and to analyze traffic. Information about your use of this site is shared with Google. By using this site, you agree to its use of cookies.

LEARN MORE GOT IT





mrxpalmeiras

- Home
- Wireshark cheat sheet
- Kail Pentest Netsec
- Rust
- ▼ Ansible
- Crystal
- ▼ Docker
- Elastic
- KDB Cheatsheet
- ▼ Git
- ▼ Golang
- JavaScript
- ^ Linux
- Iptables Basics
- Linux Cheat Sheet
- OpenVPN - Client config
- Hi Frequency/Volume Trading - OS Tuning
- SSH Certificate-based Authentication
- Troubleshooting frozen system
- RAID levels
- Barrier - screen control across physical devices
- QEMU Virtualization
- TCP troubleshooting
- Marten web framework
- Graylog
- htmx & hyperscript
- UDP packet loss troubleshooting
- SSD types and terminology
- Data Structures
- ▼ More...
- Julia cheatsheet
- Sonicwall
- Ninja Tools
- ▼ Python
- ▼ Ruby
- Ruby Cheat Sheet
- ▼ SaltStack
- Scratchpad
- ▼ Vagrant
- Vue JS Cheatsheet

Memory / Diagnostics

Debian - CPU and Mem  
lshw -html > /tmp/specs.html

Show Hardware information  
inxi -Fxzd

Fedora - CPU and Mem  
cat /proc/cpuinfo  
cat /proc/meminfo  
lspci -v

Show memory usage  
free -m

Show processes by memory usage  
ps aux | awk '{print \$6/1024 " MB\t\t" \$11}' | sort -n

show actual Memory information (RSS, memory address ,etc )  
pmap -p <PID>

show USB info  
lsusb -v

show size of folder  
du -sh

drop all memory caches  
echo 1 | sudo tee /proc/sys/vm/drop\_caches > /dev/null

Swap  
clear swap space  
swaponoff -a (wait till clears)  
swapon -a

SMEM (Memory usage profiler)

check which user/proc is using swap by %  
yum install smem  
smem -t -p -s swap

show memory usage just for my user  
smem -u

show memory usage by user  
smem -u joe

show memory usage by proc  
smem -p firefox

show memory by RSS, PSS, order by Columns  
(RSS, resident set size=portion of memory in RAM, rest in swap)  
(PSS, proportional set size=portion of main memory, RAM, occupied by proc)  
smem -c "name user pss rss"

DSTAT

show Out of Memory oom procs that are high on list to be killed  
dstat --top-oom (yum install dstat)

check process OOM score  
cat /proc/PID/oom\_adj (-10 is lower priority to get killed than 10)

Journalctl

tail a log for a process  
journalctl -u httpd -f

show last 100 lines for a process  
journalctl -u httpd --no-pager -n100

tail a process log  
journalctl -f -u <process-name>

see journal disk usage  
journalctl --disk-usage

clear journal log space anything older than 5 days  
journalctl --vacuum-time=5d

keep only last 500mb  
journalctl --vacuum-size=500M

see previous Kernel boot messages

```
journalctl --list-boots (higher numbers are older boots)
-3 89bb7913b7f84948a1dc4e05baa5c606 Tue 2023-02-14 12:58:01 CST-Tue 2023-02-14 15:16:58 CST
-2 e3605a8c134b4c1a86e4576365dddc0a Tue 2023-02-14 15:19:16 CST-Tue 2023-02-14 15:28:37 CST
-1 9d174c4a278241db85df5e38b9d17b19 Tue 2023-02-14 15:30:55 CST-Tue 2023-02-14 15:42:51 CST
```

show bootlog from boot # 3  
journalctl -b-3

show only journal logs after a certain date/time  
journalctl -S "2020-01-12 07:00:00"

DMIDecode  
show bios  
dmidecode -t bios

system info  
dmidecode -t system

chassis  
dmidecode -t chassis

memory, processor, slot  
dmidecode -t memory  
dmidecode -t processor  
dmidecode -t slot

serial #  
dmidecode -s system-serial-number

Show all current users logged in  
who  
w

Send msg to all logged-in users  
wall -n "hello"

Show all loaded modules  
lsmod

insert, remove mod  
insmod fat  
rmmod fat

Show current runlevel  
runlevel

show IRQ drivers being used  
cat /proc/interrupts

show DMA channels being used (comms between I/O ports)  
cat /proc/dma

show I/O ports being used  
cat /proc/ioports

Stress Testing  
yum install stress-ng

run stress on 2 CPUs  
stress-ng --cpu 2 --timeout 10s --metrics-brief

force Out of memory kill  
stress-ng --vm 5 --vm-bytes 95% --vm-method all --verify -t 1m -v

Stress I/O load, run 5 workers that will continually R/W to temp file  
stress-ng -d 5

Run application with memory limit  
systemd-run --user -p MemoryLimit=3G google-chrome

Kill frozen process  
Alt + PrintScreen + f

Find procs using most SWAP space

```
find /proc -maxdepth 2 -path "/proc/[0-9]*/status" -readable -exec awk -v FS=":" '{process[$1]=$2;sub(/\^[ \t]+/, "", process[$1]);} END {if(process["VmSwap"] && process["VmSwap"] != "0 kB") printf "%10s %-30s\n", process["Pid"], process["Name"], process["VmSwap"]}' '{} ' \; | awk '{print $(NF-1), $0}' | sort -hr | head | cut -d " " -f2-
```

Get top 25 Memory hogs  
ps -eo pid,user,ni,rss,vsz,cputime,lstart,etimes,time,%cpu,%mem,args --sort=-rss | head -n 25

DMESG  
check kernel actions during bootup  
dmesg -T

Top, Htop

show by memory  
top -o %MEM (hit 'c' to show full command)

This site uses cookies from Google to deliver its services and to analyze traffic. Information about your use of this site is shared with Google. By using this site, you agree to its use of cookies.



mrxpalmeiras

- Home
- Wireshark cheat sheet
- Kail Pentest Netsec
- Rust
- ▼ Ansible
  - Crystal
- ▼ Docker
  - Elastic
  - KDB Cheatsheet
- ▼ Git
- ▼ Golang
- JavaScript
- ^ Linux
  - Iptables Basics
  - Linux Cheat Sheet
  - OpenVPN - Client config
  - Hi Frequency/Volume Trading - OS Tuning
  - SSH Certificate-based Authentication
  - Troubleshooting frozen system
  - RAID levels
  - Barrier - screen control across physical devices
  - QEMU Virtualization
  - TCP troubleshooting
  - Marten web framework
  - Graylog
  - htmx & hyperscript
  - UDP packet loss troubleshooting
  - SSD types and terminology
  - Data Structures
- ▼ More...
  - Julia cheatsheet
  - Sonicwall
  - Ninja Tools
- ▼ Python
- ▼ Ruby
  - Ruby Cheat Sheet
- ▼ SaltStack
- Scratchpad
- ▼ Vagrant
  - Vue JS Cheatsheet

Network / Ifaces configuration

IP command

```
show all interfaces
ip a

show specific interface
ip addr show dev em1

assign address to interface
ip addr add 192.168.5.2 dev em1

show only active interfaces
ip link ls up

bring up an interface
ip link set dev em1 up

disable an interface
ip link set dev em1 down

rename interface w/o network restart
ip link set dev em1 down

ip link set em1 name eth1
ip link set eth1 up

delete interface
ip link delete em4

bring up an interface
ip link set em1 up

change MTU on interface
ip link set em1 mtu 9000

see all routes
ip route or route -n

get route for an IP
ip route get 192.168.1.2

delete route
ip route del 192.168.1.2

add a new route via gateway
ip route add 192.168.1.2 via 192.168.1.1 dev em1

add default route
ip route add default ia 192.168.1.1 dev em1

update route with congestion window and receive window sizing
ip route change default via 192.168.38.1 dev em1 proto static initcwnd 10 initrwnd
```

show all tunnels  
ip tunnel

NetworkManager

```
show all devices
nmcli device

start a device
nmcli device connect em1

get UUID
nmcli connection show

generate UID
uuidgen eth0

add new connection
nmcli c add connection.interface-name enp1s0 type ethernet

stop managing iface with NetworkManager
nmcli d set eth1 managed no

delete interface
nmcli dev disconnect eth1

show device information (Mac address, etc)
nmcli d show
nmcli d show eth0

connect, disconnect, status for device
nmcli d connect em1

bring up / down iface
nmcli c up eth1
nmcli c down eth1

start/stop network stack
nmcli networking off (on)
nmcli n off
nmcli n on (bounce network)

reload connections
nmcli c reload

migrate legacy ifcfg connections from /etc/sysconfig/network-scripts
nmcli c migrate em1
```

VLAN config

```
add new vlan
nmcli c add type vlan con-name bond0.252 ifname bond0.252 vlan.parent bond0 vlan.id 252

add a nested vlan
nmcli c add type vlan con-name bond0.252:5 ifname bond0.252.5 dev bond0.252 id 5
```

Network Teaming/Bonding

create network team from em1, em2 ifaces

can create teaming modes based on following:

- 1) broadcast - transmits data over all ports
- 2) roundrobin - transmits data over all ports in turn
- 3) activebackup - transmits data over one port while the other are kept as backup
- 4) loadbalance - transmits data overa ll ports with active Tx load balancing
- 5) random - random selected port
- 6) lacp - 802.3ad link aggregation protocol

- 1. create team iface  
nmcli c a type team con-name bond0 ifname bond0 team.runner loadbalance
- 2. add ifaces to team  
nmcli c a type team-slave con-name em1 ifname em1 master bond0  
nmcli c a type team-slave con-name em2 ifname em2 master bond0
- 3. configure bond0 IP details  
nmcli c m bond0 ipv4.addresses 192.168.40.20/24  
nmcli c m bond0 ipv4.gateway 192.168.40.1  
nmcli c m bond0 ipv4.dns 8.8.8.8  
nmcli c m bond0 ipv4.method manual  
  
all configs are saved into /etc/NetworkManager/system-connections
- 4. restart NM  
systemctl restart NetworkManager
- 5.

Check Traceroute and Ping at same time, live stream

mtr www.google.com

Check Port communication

find process thats holding a certain port #

netstat -tulpn | grep 5000

Netcat

Chat client

```
On Server - start NC session
hostA: nc -l 9933

on Client, connect to NC session
hostB: nc hostA 9933

can type messages between servers like chat client
```

netcat from specific interface

nc hostname 22 -s 192.168.30.23 -v

start a Netcat Bash session (ghetto SSH)

```
serverA> nc -l 5000 -e /bin/bash
serverB> nc serverA 5000
```

Netcat Ghetto web server

```
while true ; do nc -l -p 1500 -c 'echo -e "HTTP/1.1 200 OK\n\n $(date)"; done
```

Scan a range of IPs for an open port,

```
for i in {1..25};do nc -zv 208.224.251.$i 8003 -w 2 ;done
```

Scan an IP for open ports (Ghetto Nmap)

```
nc myhost 1-100 -zv
(will scan ports 1-100 and report if open or not)
```

Spin up a webserver with custom port, check that you can connect to port

```
python2 serverA> python -m SimpleHTTPServer 8331
python3 serverA> python3 -m http.server 8331
serverB> nc serverA 8331
```

connect on a UDP port

nc -u <hostname> <port> -vv

transfer files between 2 hosts

```
hostA> netcat -l 4444 > /tmp/file1
hostB> echo "cats suck dogs rule" > myfile
hostB> nc hostA 4444 < myfile
hostA> cat /tmp/file1
cats suck dogs rule
```

proxy a port via another host (similar to Socat and Redir)

A needs to connect to C:8333, but doesnt have direct access,  
A will use B as a hop to C:8333

```
hostB> nc -k -l 0.0.0.0 8333 --sh-exec "nc hostC 8333"
hostA> nc hostB 8333 -v (will conn A > B:8333 > C:8333)
```

NPING(part of nmap pkg)

send TCP packets over port 22, 80 and 443, send 500 packets at rate of 60 packets / sec with sleeptime of 3 seconds between attempts

```
(for UDP use --udp [hostname/IP]
nping --tcp nycweb01 -p 80,443,22 -c 500 -rate 60 --delay 3
```

send ARP request to all hosts on subnet  
nping -arp 192.168.30.0/24

send ICMP echo  
nping nycweb01 -icmp -icmp-type echo

send packets to ports 20-35  
nping -tcp nycweb01 -p20-35

send UDP packets  
nping -udp -c 2 -p 23000 <target>

To spin up webserver on specific network interface,

```
python -c 'import BaseHTTPServer as bhs, SimpleHTTPServer as shs;
```

Tracepath # of hops for HTTP request (better than traceroute)

```
tracepath 123.123.21.2
tracepath nycweb1
```

check link speed of iface  
ethtool em1 (speed: x)

check TCP statistics  
netstat -s -t

check congestion and other info  
netstat -s

show drop packet statistics for iface  
ip -s link show em1

query statistics for iface  
column -t /proc/net/dev

Check if port 120 is open and listening  
netstat -an | grep 120

SS (like netstat)

check user, PID listening on port 8080  
ss -ap4 | grep 8080

show all TCP connections  
ss -t

show all Listening TCP conns  
ss -lt

show all UDP connections (for Listening, add -lu)  
ss -u

display PIDs of sockets  
ss -p

filter by port number  
ss -at '( dport = :22 or sport = :22 ) '

show conns from specific source or dest address  
ss src <IP address>  
ss dst <IP address>

TCP Dump

show all interfaces tcpdump can listen on  
tcpdump -D

listen on specific interface  
tcpdump -i eth0

listen on all ifaces  
tcpdump -i any

listen on specific port or portrange  
tcpdump portrange 3334-3380  
tcpdump -i any port 12345

listen on multiple ports  
tcpdump -i any port '(80 or 443)'

search for specific src IP and port over an iface  
tcpdump port 1234 and src 1.1.1.1 -i em1

search for specific subnet  
tcpdump port 1234 and net 1.2.3.4/24

ord packet capture into a .cap file  
pdump -w capture.cap

nd contents of a .cap file  
pdump -r capture.cap

play only IP address and ports instead of hostnames  
pdump -n

play only where destination IP is 192.168.5.1 (for source use -n src)  
pdump -n dst host 192.168.5.1

capture TCP packets where port is between 1 and 1023

This site uses cookies from Google to deliver its services and to analyze traffic. Information about your use of this site is shared with Google. By using this site, you agree to its use of cookies.

LEARN MORE

GOT IT



mrxpalmeiras

- Home
- Wireshark cheat sheet
- Kail Pentest Netsec
- Rust
- Ansible
- Crystal
- Docker
- Elastic
- KDB Cheatsheet
- Git
- Golang
- JavaScript
- Linux
  - Iptables Basics
  - Linux Cheat Sheet
  - OpenVPN - Client config
  - Hi Frequency/Volume Trading - OS Tuning
  - SSH Certificate-based Authentication
  - Troubleshooting frozen system
  - RAID levels
  - Barrier - screen control across physical devices
  - QEMU Virtualization
  - TCP troubleshooting
  - Marten web framework
  - Graylog
  - htmx & hyperscript
  - UDP packet loss troubleshooting
- SSD types and terminology
- Data Structures
- More...
  - Julia cheatsheet
  - Sonicwall
  - Ninja Tools
- Python
- Ruby
  - Ruby Cheat Sheet
- SaltStack
- Scratchpad
- Vagrant
  - Vue JS Cheatsheet

```
tcpdump -n tcp dst portrange 1-1023

capture packets where destination host is 192.168.5.1 and port is 5049
tcpdump -n "dst host 192.168.5.1 and dst port 5049"

print packets in ASCII or HEX
tcpdump -i any -A (ascii) or -X (hex)

run in background and record to file
nohup tcpdump -i any port 27025 -w myfile.cap &

rotate pcap file similar to logrotate, this will create 10 pcap files of 200MB each
tcpdump -n -W 10 -C 200 -w /tmp/file.pcap

check multicast data
tcpdump -i p1p2 -s0 -vvv host 233.143.214.1
```

Flags

```
[.] - ACK (Acknowledgment)
[S] - SYN (Start Connection)
[P] - PSH (Push Data)
[F] - FIN (Finish Connection)
[R] - RST (Reset Connection)
[S.] - SYN-ACK (SynAck Packet)
```

Wireshark

show only bad packets  
add to filter: tcp.analysis.flags

show only Resets  
tcp.flags.reset == 1

show packets between dates  
(frame.time >= "July 27, 2020 08:40:00" && frame.time <= "July 30, 2020 08:42:42")  
show only problematic packets  
\_ws.expert.severity == error

show only UDP / TCP/ RTCP  
type into filter: udp (or tcp, rtcp)

Check MAC address mapping to IP  
arp

clear ARP cache  
ip -s -s neigh flush all

IP Routing Table  
route -n

add a new route,  
ip route add 118.100.1.173 via 192.168.38.17 dev p1p2 metric 200

add new route permanently  
vim /etc/sysconfig/network-scripts/route-p1p2  
201.224.250.40 via 192.168.38.33 metric 200

delete a route  
ip route del 118.100.1.173

modify existing route  
ip route del 40.2.2.0/24 via 30.1.2.2  
ip route add 40.2.2.0/24 via 30.1.2.2 metric 1234  
# kill all connections on port 21  
tcpkill -i eth0 port 21

add TCP permissions to TCP analyzer tools so non-root users can create sockets and access network interfaces  
setcap cap\_net\_raw,cap\_net\_admin=eip /usr/bin/tcppreplay  
setcap cap\_net\_raw,cap\_net\_admin=eip /usr/bin/tcpdump

IPERF  
check bandwidth usage  
yum install iperf3

on server:  
iperf3 -s

on client:

```
iperf3 -c <IP of server> -p 5001 <port> -P 20 <# of parallel TCP conns> -t 20 <run for x seconds>

Client connecting to 208.224.251.3, TCP port 5001
TCP window size: 90.0 KByte (default)
-----
[ 3] local 172.31.123.96 port 48908 connected with 208.224.251.3 port 5001
write failed: Connection reset by peer
[ ID] Interval      Transfer    Bandwidth
[ 3] 0.0- 0.0 sec    130 KBytes  68.9 Mbits/sec
```

test bw using 10 parallel sessions, each session sending maximum of 2mb  
iperf3 -P 10 -b 2M -c <hostname>  
  
use UDP iperf, will flood the connection with datagrams, without session limits of TCP, more accurate bandwidth read than TCP  
iperf3 -u -b 2M -c <hostname>  
  
show iperf3 details (congestion, window sizes, etc)  
iperf3 -c nycweb1 -t 15 --debug

Nuttcp - advanced iperf

check bandwidth usage on a host directly  
iftop -PN

SAR

show live statistics of traffic over all interfaces  
sar -n DEV 1 2  
DEV = network iface info  
EDEV = network errors  
NFS = active NFS clients  
NFSD = NFS server info  
SOCK = socket info  
ALL = all above

Parameter Description:

IFACE: LAN interface  
rxpck/s: packets received per second  
txpck/s: packets sent every second  
rxbyt/s: number of bytes received per second  
txbyt/s: number of bytes sent per second  
rxcmp/s: compressed packets received per second  
txcmp/s: compressed packets sent every second  
rxmcst/s: multicast packets received per second  
rxerr/s: bad packets received per second  
txerr/s: bad packets sent every second  
coll/s: conflicts per second  
rxdrop/s: the number of received packets dropped per second because the buffer is full  
txdrop/s: the number of sent packets dropped per second because the buffer is full  
txcarr/s: number of carrier errors per second when sending packets  
rxfram/s: the number of frame alignment errors received per second  
rxfifo/s: the number of FIFO over speed errors per second of received packets  
txfifo/s: the number of FIFO over speed errors per second in packets sent

show statistics for all ifaces  
cat /proc/net/dev

```
bhs.HTTPServer(("192.168.200.99", 8331),
shs.SimpleHTTPRequestHandler).serve_forever()'
```

check ports using nmap

nmap	localhost
PORT	STATE SERVICE
22/tcp	open ssh
25/tcp	open smtp
80/tcp	open http
89/tcp	open su-mit-tg

NMAP

check subnet for open ports  
nmap -sP -PS22,3389 192.168.30.1/24

DNS

Check DNS routing

```
host github.com
github.com has address 192.30.253.113
github.com has address 192.30.253.112
github.com mail is handled by 10 ALT3.ASPMX.L.GOOGLE.com.
```

Dig into DNS query  
dig www.domain.com

check all DNS name servers  
cat /etc/resolv.conf

get your public IP from google

dig +short myip.opendns.com @resolver1.opendns.com  
124.245.66.135

or this  
curl -4 icanhazip.com

Check all open network connections  
lsof -i

Check which procs are holding up deleted files  
lsof +L1

check output of df vs du  
df shows total usage including file descriptors, du shows actual usage  
kill any procs holding up "deleted" file descriptors, will show reduction of used space

Get true Timezone  
curl https://ipapi.co/timezone

Multicast

see what MC groups are present  
ip maddr

Network Utilities

hping3 - like ping but can connect to ports and use TCP  
iftop - iface network activity top  
ss - better version of netstat (ss -ap4)  
iptraf - interface and network cmd line gui tool (very good)

Kill a TCP session w/o killing process (will only kill new connections, not Established)  
yum install dsniff  
tcpkill -i eth0 port 28394

kill Established TCP connection via port (doesnt kill parent process)

lsof -np <PID of Parent> | grep <IP of remote host> (get the FD number, 4th column)  
gdb -p <PID of parent> --batch -ex 'call shutdown(FD #)'

ie, need to kill this specific TCP session but not kill MyApp (this app has other TCP established connections)  
tcp 0 0 192.168.38.21:25959 108.124.250.173:50443 ESTABLISHED 221955/MyAPP

lsof -np 221955 | grep 108.124.250.173  
risk\_gate 221955 qbsim 17u IPv4 2857516568 0t0 TCP 192.168.38.21:25959->108.124.250.173:51212 (ESTABLISHED)

FD id = 17u (update), now free up this file descriptor  
gdb -p 211955 --batch -ex 'call shutdown(17u, 2)'

Close File descriptor without killing the process (ie proc is up but file is deleted)  
check for deleted files

lsof +L | grep deleted | grep <filename> (get PID of this proc)

get ID of FD (4th column of lsof output), ie 43w

detach file descriptor from proc

gdb -p <PID> --batch -ex 'p close(43)'

---

This site uses cookies from Google to deliver its services and to analyze traffic. Information about your use of this site is shared with Google. By using this site, you agree to its use of cookies.

LEARN MORE GOT IT





mrxpalmeiras

- Home
- Wireshark cheat sheet
- Kali Pentest Netsec
- Rust
- Ansible
- Crystal
- Docker
- Elastic
- KDB Cheatsheet
- Git
- Golang
- JavaScript
- Linux
  - Iptables Basics
  - Linux Cheat Sheet
  - OpenVPN - Client config
  - Hi Frequency/Volume Trading - OS Tuning
  - SSH Certificate-based Authentication
  - Troubleshooting frozen system
  - RAID levels
  - Barrier - screen control across physical devices
  - QEMU Virtualization
  - TCP troubleshooting
  - Marten web framework
  - Graylog
  - htmx & hyperscript
  - UDP packet loss troubleshooting
- SSD types and terminology
- Data Structures
- More...
  - Julia cheatsheet
  - Sonicwall
  - Ninja Tools
- Python
- Ruby
  - Ruby Cheat Sheet
- SaltStack
- Scratchpad
- Vagrant
- Vue JS Cheatsheet

IPTables

show all rules

iptables -L -n -v

show all FORWARD rules

iptables -L FORWARD --line-numbers

delete a rule

iptables -D FORWARD <line number>

check existing NAT rules

iptables -t nat -v -L POSTROUTING --line-number

iptables -t nat -v -L PREROUTING --line-number

forward any request from ServerA port 80 to ServerB port 80 on server A

(make sure to add sysctl -w net.ipv4.ip\_forward=1)

iptables -t nat -A PREROUTING -p tcp --dport 80 -j DNAT --to-destination <IP of serverB>:80

iptables -t nat -A POSTROUTING -p tcp -j MASQUERADE

change outgoing packets IP header

iptables -t nat -A POSTROUTING -d <destination IP> --dport <PORT> -j SNAT --to-source <IP you want to change to>

forward an OUTGOING packet for a specific port (going from host A), to another host (host B) host A>

iptables -t nat -A OUTPUT -p tcp --dport 8331 -j DNAT --to-destination 10.182.26.8:8331

allow all connections from an IP

iptables -A INPUT -s 59.50.131.179 -j ACCEPT

forward a packet going to a specific hostname and port to another hostname:port

iptables -t nat -A PREROUTING -p tcp -d 18.224.251.4 --dport 22 -j DNAT --to-destination 192.168.10.22:22

completely flush all chains, rules, filter, raw, mangle, etc  
## allow all incoming connections to avoid being locked out  
iptables -F INPUT ACCEPT  
# flush custom chains, nat, raw, security, mangle, filter rules  
iptables -X  
iptables -t nat -F  
iptables -t raw -F  
iptables -t security -F  
iptables -t mangle -F  
iptables -F

Packages / Libs / Modules

show installed software

Debian distro

dpkg -l

apt-cache search [pkg name]

Fedora distro

yum list installed

rpm -qa | grep [pkg name]

yum search [pkg name]

RPM install package

rpm -i pkg.rpm

rpm -i mypkg.rpm --force (force install)

rpm -i mypkg.rpm --nodeps (ignore dependencies)

what RPM does a file belong to?

rpm -qf /usr/bin/mysqlaccess

show files inside installed RPM package

rpm -ql package-name

show files inside local uninstalled RPM package

rpm -qpl local-file.rpm

Show libraries for a program

ldd /bin/ls

refresh YUM cache

yum clean expire-cache|

yum clean all

show dependency for a package

yum -q deplist \$pkg

see install/upgrade history

yum history

get info on specific yum transaction

yum history info <# of transaction>

rollback yum patch

yum history undo <# of transaction>

save all IPTABLES rules permanently

iptables-save > /etc/sysconfig/iptables

restore from file

iptables-restore < /tmp/backup.iptables

add Debug log to prerouting rule #3 (tail syslog)

iptables -t nat -I PREROUTING 3 -j LOG

allow SSH port 22 only from address 190.120.30.3, block all others

iptables -I INPUT -p tcp '!' -s 190.120.30.3 --dport 22 -j REJECT

allow SSH port for specific address

iptables -A INPUT -p tcp -s 190.120.30.3 --dport 22 -j ACCEPT

block port

iptables -A OUTPUT -p tcp --dport 2500 -j DROP

allow a port

iptables -A INPUT -p tcp --dport 2500 -j ACCEPT

allow an IP address

iptables -A INPUT -p tcp -s 192.168.3.5 -j ACCEPT

iptables -A OUTPUT -p tcp -d 192.168.3.5 -j ACCEPT

block an IP address

iptables -A INPUT -s 192.130.2.4 -j DROP

block range of IPs

iptables -A INPUT -s 192.168.2.0/24 -j DROP

allow range of ports (1200 and 5000-6000)

iptables -A INPUT -p tcp --match multiport --dports 1200,5000:6000 -m conntrack -j ACCEPT

redirect port to another port on same host

iptables -t nat -A PREROUTING -i eth0 -p tcp --dport 25 -j REDIRECT --to-port 2525

create a custom CHAIN

iptables -N My-Custom-Rules

rebuild cached library list or add new libs

vi /etc/ld.so.conf

ldconfig

Install package including its dependencies, example 'mysql'

yum deplist mysql | awk '/provider:/ {print \$2}' | sort -u | xargs yum -y install

show installed packages by disk space usage (Centos)

rpm -qa --queryformat "%10(size) - %-25(name) %t %%(version)\n" | sort -n

remove old kernels, keep only current and 1 older

package-cleanup --oldkernels --count=2

Modules

show custom modules

dkms status (yum install dkms)

show loaded modules

lsmod | grep <modname>

load module (insert)

insmod /lib/modules/<kernel version>/kernel/drivers/<etc>

unload module

rmmod /lib/modules/<kernel version>/kernel/drivers/etc

same but using modprobe w/o needing path to modules

insert/load: modprobe <modname>

remove: modprobe -r <modname>

This site uses cookies from Google to deliver its services and to analyze traffic. Information about your use of this site is shared with Google. By using this site, you agree to its use of cookies.

LEARN MORE GOT IT



- Home
- Wireshark cheat sheet
- Kail Pentest Netsec
- Rust
- Ansible
- Crystal
- Docker
- Elastic
- KDB Cheatsheet
- Git
- Golang
- JavaScript
- Linux
  - Iptables Basics
  - Linux Cheat Sheet
  - OpenVPN - Client config
  - Hi Frequency/Volume Trading - OS Tuning
  - SSH Certificate-based Authentication
  - Troubleshooting frozen system
  - RAID levels
  - Barrier - screen control across physical devices
  - QEMU Virtualization
  - TCP troubleshooting
  - Marten web framework
  - Graylog
  - htmx & hyperscript
  - UDP packet loss troubleshooting
- SSD types and terminology
- Data Structures
- More...
  - Julia cheatsheet
  - Sonicwall
  - Ninja Tools
- Python
- Ruby
  - Ruby Cheat Sheet
- SaltStack
- Scratchpad
- Vagrant
  - Vue JS Cheatsheet

## Process / Init / CPU

```
get uptime of a process
ps -p $$ -o etime=
where $$ is PID, result is in format dd-hh:mm:ss

find PID of a process (add to .bashrc)

function pid() { ps -fu $USER | grep $1 | grep -v "grep" | grep -v "ps -fu"  ;}

Run process in background (use & to push to background)

./run_script.sh &

Get current PID
$$

kill all processes by name (with confirmation)
pkill -f $name

kill process by Port
fuser -k 5100/tcp

kill process by owner name
killall -u username

find process by name, kill all
ps -ef | grep "vault server" | grep -v grep | awk '{print $2}' | xargs kill -9

show all procs and their children
pstree -ap

show 4 way scrollable process tree
ps awwfix | less -S

show all processes and children
ps -ef --forest

show # of processes per user
ps hax -o user | sort | uniq -c | sort -r

kill a process running on Port 8331
kill -9 $(lsof -i :8331 | awk '{l=$2} END {print l}')
```

```
get amount of open file descriptors by user
lsof -u <username> | wc -l

get home directory of a process ID
pwdx <PID>
```

### NTP

```
check offset of time between 2 servers,
[23:38 root@web1:~]# ntpdate -q web2
server 10.112.42.8, stratum 2, offset 0.005212, delay 0.02580
13 Aug 23:39:01 ntpdate[17325]: adjust time server 10.182.48.8 offset 0.005212 sec
```

offset of less than 5/1000s of a second

```
check offset against timeserver
ntpq -p
```

```
run in debug
ntpdate -dv <name of timeserver>
```

```
Synchronize time w another host over SSH (server2 has correct date)
date --set="$(ssh user@server2 date)"
```

```
Hide processes and PIDs for non-root users
edit /etc/fstab
proc /proc proc defaults,hidepid=2 0 0

remount
mount -o remount,rw,hidepid=2 /proc
```

```
to add an exception for a group/user (let this group see other PIDs), add 'gid' & remount
proc /proc proc defaults,hidepid=2,gid=joe 0 0
```

```
Isolate CPUs for specific processses
grubby --default-kernel
/boot/vmlinuz-3.10.0-862.14.4.el7.x86_64

grubby --info=/boot/vmlinuz-3.10.0-862.14.4.el7.x86_64
args="ro no_timer_check console=tty0 console=ttyS0,115200n8
net.ifnames=0 biosdevname=0 elevator=noop crashkernel=auto
LANG=en_US.UTF-8"

## get current isolated cores
cat /sys/devices/system/cpu/isolated

## add cpu isolation
grubby --update-kernel=/boot/vmlinuz-3.10.0-862.14.4.el7.x86_64 --
args=isolcpus=2,3

reboot host to pickup changes

## remove isolation
grubby --remove-args="isolcpus=2,3" --update-kernel=<kernel name>
```

## User / Group / Sudo

### USERS

```
create new user
adduser eric

add user to Group
usermod -aG mygroup eric

remove user from Group
gpasswd -d <user> <group>

add user to multiple groups
usermod -aG group1,group2,group3 eric

change UID for user
usermod -u 2550 eric

change GID for user
groupmod -g 2550 eric

lock a user account
passwd -l eric

unlock user account
passwd -uf eric

delete a user's password
passwd --delete eric

change user's shell
usermod --shell /bin/bash eric

remove expired password requirements for user
chage -m 0 -M 99999 -I -1 -E -1 jsmith

remove expired password for all users
for user in $(cut -d: -f1 /etc/passwd); do sudo chage -M 99999 -I -1 -m 0 -E -1 $user; done

delete /home folders of users that are deleted on system
(check for user GroupID, if 'UNKNOWN', then delete that home folder)
for i in $(ls /home) ; do stat --format='%G' $i | xargs echo $i | grep UNKNOWN |
awk '{print "/home/"$1}' | xargs rm -rf;done
```

Run command with a process "niceness" or priority (-20 highest priority, 19 lowest)

```
nice -18 cat /etc/hosts

Check new incoming connections on port, live
ss -nap | grep 4433
```

```
Change a running program's priority (change to priority 7, PID 168390 for all processes running by
users 'root' and 'joe'
renice 10 168390 -u root joe
```

### Systemctl

```
show all enabled services
systemctl list-unit-files | grep enabled
```

```
show all running services
systemctl list-units --type=service --state=running
```

```
start / stop / status / refresh / reload / enable / disable / show
systemctl start httpd.service
```

```
analyze bad startup script
systemd-analyze verify monit.service
```

```
refresh sysctl
systemctl daemon-reload
```

```
I/O
monitor high disk IO
* * * * * root /usr/sbin/iotop -botqqqk --iter=60 | grep -P "\d\d.\d\d K/s" >>
/var/log/iotop
```

```
Limit CPU usage for a process #2240 to 50% of CPU and also its child procs
cpulimit -pid 2240 -l 50 -i
```

```
Taskset and NUMACTL
start a process on only 1st CPU core
taskset -c 0 /bin/nginx
```

```
for multiple CPU affinity
nohup taskset -c 0,1,2,5 /bin/program
```

```
get range of CPUs on which process can run on (affinity)
taskset -cp <PID>
```

```
get CPU on which a PID is running on
ps -mo pid,tid,fname,user,psr -p <PID>
```

```
pin processes to specific CPUs that are isolated (by default, numa does not allow pin to isolated
CPUs, must use ALL option)
cat /proc/cmdline
isolcpus=2,3
nohup numactl --all -C 2,3 /bin/myprogram
```

```
find User and Parent PID of a zombie process thats holding up a port
#1 get the iNODE
root@min1# netstat -ltnae | awk 'NR==2 || /:18100/'
Proto Recv-Q Send-Q Local Address          Foreign Address         State       User    Inode  PID/Program name
tcp        1      0 0.0.0.0:18100          0.0.0.0:*               LISTEN      1000    24444060 -
tcp        1      0 192.168.37.5:18100    208.224.250.11:1046    CLOSE_WAIT  0       0      -

#2 search by iNODE
root@min1# lsof | awk 'NR==1 || /24444060/'
COMMAND  PID  TID     USER  FD  TYPE             DEVICE  SIZE/OFF      NODE NAME
trading_ 1385 17 138529   joe   50u  IPv4             24444060  000      TCP *:18100 (LISTEN)
```

```
run program in background, no output
nohup programName 2>&1 &
```

```
clear abrt'd messages
if getting abrt'd-cli timed out
check /var/spool/abrt or /usr/local/spool/abrt
remove old abrt files, restart abrt'd service
```

```
Generate Core Dump file into specific location
sysctl -w kernel.core_pattern = /mnt/core.%e.%p.%h.%t
```

```
add core limits, set limit of core size to 5mb - 4096 bytes per block
5(MB) * 1024 * 1024 / 4096
1mb = 256 blk
1gb = 262,144 blk
```

```
vim /etc/security/limits.conf
joe soft   core    1280 (4096 bytes per block, 5MB core = 1280 blocks)
joe hard   core    1280
```

```
check core limits (start new session as Joe)
joe> ulimit -a
create new background proc
joe> nohup python -m SimpleHTTPServer &
kill proc to generate core file
joe> kill -s SIGTRAP $(pgrep python)
```

```
remove user from group
gpasswd -d joe wheel
```

```
create a nologin user (no home dir)
useradd -r joe
or
adduser -r -s /bin/nologin jsmith
```

```
create user Joe with custom home dir, custom ID 999, custome group ID 555, add to 2 groups
 corp, webadmins)
useradd -d /var/home/joe -u 999 -g 555 -G corp,webadmin joe
```

```
via Perl
adduser --home /var/home/joe -u 999 -g 555 -G corp,web joe
```

### GROUPS

```
add new group
groupadd mygroup
```

```
remove group
groupdel mygroup
```

```
modify group ID
groupmod -g 999 mygroup
```

```
change group name
groupmod -n newgroup oldgroup
```

```
show what cores a process is running on
for i in $(pgrep <name of process>); do ps -mo pid,tid,fname,user,psr -p $i;done
```

```
allow user to run command as another user (joe can run httpd command as fred)
vi /etc/sudoers.d/httpd
joe ALL=(fred) NOPASSWD: /bin/httpd

joe> sudo -u fred /bin/httpd
```

```
or another way
joe> su -c 'bash myprogram.sh' fred
```

```
Check sudoers syntax
visudo -cf /etc/sudoers.d/mysudo
```

This site uses cookies from Google to deliver its services and to analyze traffic. Information about your use of this site is shared with Google. By using this site, you agree to its use of cookies.





mrxpalmeiras

- Home
- Wireshark cheat sheet
- Kail Pentest Netsec
- Rust
- ▼ Ansible
  - Crystal
- ▼ Docker
  - Elastic
  - KDB Cheatsheet
- ▼ Git
- ▼ Golang
- JavaScript
- ^ Linux
  - Iptables Basics
  - Linux Cheat Sheet**
  - OpenVPN - Client config
  - Hi Frequency/Volume Trading - OS Tuning
  - SSH Certificate-based Authentication
  - Troubleshooting frozen system
  - RAID levels
  - Barrier - screen control across physical devices
  - QEMU Virtualization
  - TCP troubleshooting
  - Marten web framework
  - Graylog
  - htmx & hyperscript
  - UDP packet loss troubleshooting
- SSD types and terminology
- Data Structures
- ▼ More...
  - Julia cheatsheet
  - Sonicwall
  - Ninja Tools
- ▼ Python
- ▼ Ruby
  - Ruby Cheat Sheet
- ▼ SaltStack
- Scratchpad
- ▼ Vagrant
  - Vue JS Cheatsheet

File / Dir

**Rsync**  
sync files from one Dir1 to Dir2  
rsync -azP dir1/ dir2    ## -z flag is compression

*-azP flag is used to compress file (z), and P for partial, it will only rsync deltas instead of starting all over from scratch*

RSYNC file to a remote system's /tmp dir  
rsync -azP file1 root@remotesystem:/tmp

rsync and exclude logs, png  
rsync -azP --exclude={\*.log,\*.png} server1:/tmp/dir /tmp

Pull file from a remote system to a local /tmp dir  
rsync -azP root@remotesystem:/opt/file1 /tmp

If Rsync not found, use path  
  
--rsync-path=/usr/bin/rsync

**Rsync using a hop server (A > B > C)**  
assuming you can ssh joe@A > joe@B  
and can ssh from joe@B > joe@C

rsync -azP -e "ssh -A joe@B ssh" file1 joe@C:/tmp

will rsync local file1 via B, into C

if Rsync versions dont match up, can also do this, (rsyncs file on C to localhost via B)  
rsync -azP -e 'ssh -o "ProxyCommand ssh -A joe@B nc %h %p"' joe@C:/tmp/xferfile .

rsync - set mod and ownership on incoming files/dirs,

```
hostA> ls -la /home/joe
drwxrwsr-x. 3 joe groupA  21 Sep 16  2018 tmp/

hostB> rsync -azP --chmod 644 --chown=mary:accounting hostA:/home/joe/tmp .
hostB> ls -la
drw-r--r--  mary  accounting  /tmp
```

Rsync using specific SSH keys

**Sort file**  
sort -d filename ## alphabetically  
sort -r filename ## reverse order  
sort -n filename ## numeric sort  
sort -M filename ## sort by month date

SSHFS  
sudo sshfs -o allow\_other,defer\_permissions root@xxx.xxx.xxx.xxx:/mnt/droplet

copy all files to destination except for whatever is in .gitignore  
cp -r !({cat .gitignore}) /tmp/dest

mount NFS share  
yum install nfs-utils nfs-utils-lib  
service nfs start  
mount -t nfs <serverIP>:/path/of/mount /mnt/point

remove first 500 lines of a file, in place (shrink a log file)  
sed -i -e 1,500d file.log

reduce log file to 200b  
truncate -s 200 file.log

**User & Group Permissions**  
give 'sysadmin' Group 777 permission to a dir /opt/test  
chmod g+rxw /opt/test

change group ownership for symlink (recurse down)  
chgrp -Rh mygroup /home/user/dir

add execute bit for group on all folders  
find . -type -d | xargs chmod g+x

change group ownership of a dir  
chgrp sysadmins /opt/test

Get ACL on a directory  
getfacl /opt/test

give Sysadmins group 777 to /opt/test  
setfacl -m group:sysadmins:rwX /opt/test

to set recursively down,  
setfacl -Rm u:joe:rwX /home/mary

remove ACL  
setfacl -x user:antony /opt/test

give r/w access to /home/user1 and preserve SSH security  
chmod 750 /home/user1  
setfacl -m user:user2:rw /home/user1

remove all ACLs from file or dir  
setfacl -b /home/user1

set a default ACL for a directory (all new files or dirs created in this directory will inherit ACL permissions)  
setfacl -d -m u::rwx,g::rwx,o::- /opt/testdir  
setfacl -Rdm u:joe:rwx /opt/somedir

Backup and restore all permissions  
make a backup of all permissions in a directory,  
getfacl -R /home/user > /tmp/permissions\_backup

restore all perms recursively  
setfacl --restore=/tmp/permissions\_backup

ensure all files and dirs created by user, inherit the Group permission of parent directory (SUID bit)  
- this example gives Joe rwx, gives group "employees" only Read (directories get set with X in order for group members to 'ls' to them), all others have no access to this folder or subfolders

1. chgrp -Rh employees /home/joe
2. setfacl -d -Rm u::rwx,g::rwx,o::- /home/joe
3. chmod -R g+s /home/joe (set S bit to inherit parent permissions for all new subfolders)
4. chmod -R g-w /home/joe (removes write perms for group inside joe's home folder)

add timestamp to a tail of log file  
tail -f /var/log/messages | while read ; do echo "\$(date +%T.%N) \$REPLY" ; done

copy all ssh keys for every user from 1 host to another  
host1> for i in \$(ls /home);do rsync -azP /home/\$i/.ssh/id\_rsa\*  
host2:/home/\$i/.ssh/ ;done

get directory permissions of a user's directory in numeric form  
stat -c '%a' /home/user    >> 700  
  
get owner of directory  
stat -c '%u' /home/user    >> user    (use %g for group)

Logrotate

place all logrotate confs in /etc/logrotate.d  
/var/log/httpd/\*log {  
  rotate 3 # how many rotated files to keep left over  
  size 10MB # rotate if log exceeds this  
  daily # rotate on daily basis unless size max criteria is met first  
  maxage 20 # delete old rotate files over 20 days  
  compress # gzip compress rotated files  
  missingok  
  notifempty  
  sharedscripts  
  postrotate  
    /sbin/service httpd graceful 1>/dev/null 2>&1 || true  
  endscript  
}

**Searching**  
find all files larger than 100M  
find /home -xdev -type f -size +100M | xargs du -sh | sort -hr  
  
Find 10 largest files  
find . -type f -print0 | xargs -0 du | sort -n | tail -10 | cut -f2 | xargs -l{} du -sh {}  
  
another way  
find /home -type f -exec du -Sh {} + | sort -rh | head -n 5

find all files created in last 120 minutes  
find / -cmin 120

Find 10 largest dirs  
find . -type d -print0 | xargs -0 du | sort -n | tail -10 | cut -f2 | xargs -l{} du -sh {}

find 25 largest files in current dir and its subdirs  
find . -type f -exec ls -al {} \; | sort -nr -k5 | head -n 25

find duplicate files, (based on MD5 hash)  
find -type f -exec md5sum '{}' ';' | sort | uniq --all-repeated=separate -w 33

find specific user's files  
find . -user <username> -print

find total size of files matching a patter  
du -ch \*.jpg | grep total

recursively remove all empty subdirs  
find . -depth -type d -empty -exec rmdir {} \;

find all hard links to a file  
find /path/to/dir -xdev -samefile <name of file>

find the latest modified files (recursively)  
find . -type f -exec stat --format '%Y :%y %n' "{}" \; | sort -nr | cut -d: -f2- | head

find files modified or created in last 2 days  
find /dir -newermt "2 days ago" -ls

Show top 10 largest open files  
lsdf / | awk '{ if(\$7 > 1048576) print \$7/1048576 "MB" " " "\$9 " " \$1 }' | sort -n -u | tail

show 10 largest files in a directory  
du -a /opt/blah | sort -n -r | head -n 10

list by size(-S), human readable(-h), all(-a), reverse date order (-r), list (-l), date (-t)

find files older than 300 days, display them  
find /tmp -type f -mtime +300 -print | xargs ls -lha

now delete them  
find /tmp -type f -mtime +300 -print | xargs rm

Find and Search  
find -name filename ## any file

Find recursively any hidden file  
find /dirname -name ".\*" -print

show only hidden files and directories  
ls -l -d .[!]?\*

Find in specific dir  
find /tmp -name myfile

Find file in specific location larger than 20MB  
find /tmp -size +20M

Find files larger than 20MB and older than 360 days, delete them  
find /tmp -type f -size +20M -mtime +300 -print | xargs rm

get last element  
echo /my/dir/name/backups/someFile.tar | awk -F"/" '{print \$(NF)}'  
someFile.tar

get filename from a base path,  
basename /my/dir/name/backups/someFile.tar // someFile.tar

compare contents of 2 directories  
diff <(cd </path/to/dir1> && find | sort) <(cd </path/to/dir2> && find | sort)

Freeze (lock) a directory or file from being modified (ACL, permissions, ownership, etc) - only root can unlock this. NOTE - this also prevents creating new files, this "freezes" the dir completely.  
chattr +i <dir name> (locks dir)  
chattr -i <dir name> (unlocks)

**filesystem shows 100% usage, but actual usage is much less (FS has too many inodes open)**  
1. check amount of free inodes on mount (ie, /home shows 100% usage)  
df -i (check INode column)

2. see which files are from dead procs

lsdf +L1 | grep /home | awk '{ \$7=\$7/1048576 " MB"}'1'

3. unmount directory

umount /home

4. repair FS (check which block device with lsblk)

xfs\_repair /dev/sda1

5. remount (mount -a)

This site uses cookies from Google to deliver its services and to analyze traffic. Information about your use of this site is shared with Google. By using this site, you agree to its use of cookies.

LEARN MORE   GOT IT



mrxpalmeiras

- Home
- Wireshark cheat sheet
- Kail Pentest Netsec
- Rust
- Ansible
- Crystal
- Docker
- Elastic
- KDB Cheatsheet
- Git
- Golang
- JavaScript
- Linux
- Iptables Basics
- Linux Cheat Sheet
- OpenVPN - Client config
- Hi Frequency/Volume Trading - OS Tuning
- SSH Certificate-based Authentication
- Troubleshooting frozen system
- RAID levels
- Barrier - screen control across physical devices
- QEMU Virtualization
- TCP troubleshooting
- Marten web framework
- Graylog
- htmx & hyperscript
- UDP packet loss troubleshooting
- SSD types and terminology
- Data Structures
- More...
- Julia cheatsheet
- Sonicwall
- Ninja Tools
- Python
- Ruby
- Ruby Cheat Sheet
- SaltStack
- Scratchpad
- Vagrant
- Vue JS Cheatsheet

Compression / Cron / Mount / Encryption

mount ISO

mount -t iso9660 -o loop /home/tecmint/Fedora-18-i386-DVD.iso /mnt/iso/

umount ISO

umount /mnt/iso

Compression

compress using bz2

tar cvfj mydir.tar.bz2 /home/mydir

untar tar.bz2 file

tar -xvf file.tar.bz2

uncompress bz2 file

bzip2 -dk file.bz2

compress a file XZ format ( best compression )

tar -cvpJf mydir.tar.xz /home/user/mydir

compress a folder into XZ format, perserve permissions, dont include parent folders, just include the target folder (folder is located in /mnt/hc/myFolder)

tar -cJf myFolder.tar.xz -C /mnt/hc myFolder --preserve-permissions

untar a XZ tarball

tar -xf myFolder.tar.xz

untar XZ to a specific directory, preserve permissions inside tarball

tar -xf myFolder.tar.xz -C /home/joe --preserve-permissions

uncompress XZ file

unxz file.xz

tar a file or dir into tar.gz

tar zcvf name.tar.gz file1 dir1 dir2

untar and unzip

tar -xvzf file.tar.gz

untar .tgz

tar xzvf file.tgz

compress using LZMA

tar -cavvf file.tar.lzma file

uncompress LZMA to a directory

tar -xavvf file.tar.lzma -C <dir>

see whats inside a tar

tar -tvf mydir.tar

see inside tar.gz

tar -tf file.tar.gz

see inside tar.xz

tar tvfa file.tar.xz

untar single file from tar.gz (for bz2, replace tar.gz with tar.bz2)

tar --extract --file=mydir.tar.gz file1

untar multiple files using wildcard

tar -zxvf mydirs.tar.gz --wildcards '\*.php'

create a symlink

ln -s <path to actual binary> <target location>

ln -s /usr/local/bin/myapp /usr/bin/myapp

download entire website down to local level (and convert links to local) Wget Mirror

wget -mk www.google.com

download a file using curl

curl -O -u<USERNAME>:<API\_KEY> -X GET

https://api.bintray.com/packages/orgname/repo\_name/pkg\_name/logs/downloads-03-12-2016.csv.gz

Make disk backups/images with dd

create a disk backup of disk /dev/sda, save to backup.gz

dd if=/dev/sda conv=sync,noerror bs=128K | gzip -c > /mnt/hc/backup.gz

restore image onto disk

gunzip -c /mnt/hc/backup.gz | dd of=/dev/sda

create backup image of host1, and store img on remote host2.

host1> dd if=/dev/sda conv=sync,noerron bs=128K status=progress | gzip -c | ssh user@host2 'dd of=/opt/backup.gz'

restore host1 by reading backup from host2

host1> ssh user@host2 'dd if=/opt/backup.gz status=progress' | gunzip -c | dd of=/dev/sda

Audit

show status of audit system

auditctl -s

show all audit rules

auditctl -l

clear all rules

auditctl -D

monitor file for any changes

auditctl -w /etc/filename -p wa -k myfile\_changes

see any changes done to file

ausearch -k myfile\_changes

save audit rules permanently

add to /etc/audit/rules.d/audit.rules

-w /etc/filename -p wa -k myfile\_changs

check user actions by user name, from yesterday to now,

ausearch -ua joe -ts yesterday -te now -i

search by specific command and specific directory that was deleted

ausearch -f /tmp/testdir -x /bin/rm

search by type of event

ausearch -ua joe -m SYSCALL (or EXECVE)

search by time range

ausearch -ua joe --start 09/09/2019 '12:04:00' --end 09/12/2019 '12:22:00'

search by parsing a specific log file

ausearch -ua joe --input /tmp/audit.log

search raw text grep by port number

ausearch -r | grep -E "a[1-5]="\9999\""

>> type=EXECVE msg=audit(1677000183.185:1366): argc=4 a0="nc" a1="titan" a2="9999" a3="-v"

get list of failed login attempts by user and IP where theyre coming from

last -f /var/log/btmp

FTP / LFTP

ftp a file providing username + password

lftp sftp://"user:password@host" -e "put -O path/on/target /tmp/file.txt"

LFTP using SSH options, run command

lftp sftp://\$user:\$ftpcred@\$host:\$port -e "set sftp:connect-program 'ssh -oHostKeyAlgorithms=ssh-rsa'; cd \$rem\_logdir; put \$logname; bye"

Disk / Partitioning / FileSystem

unmount volume

umount /mnt/nas1

kernel is holding up NFS mount on a bad connection

ount -l /mnt/nas1

check what proc is holding up unmounting

ser -mv /mnt/nas1

of | grep /mnt/nas1

all partitions

isk -l

Disk /dev/sda: 11.3 GB, 11286446080 bytes, 22043840 sectors

Units = sectors of 1 \* 512 = 512 bytes

check JSON formats for multiple files

install jsonlint and check format

npm install jsonlint -g

\$ for i in \$(ls | grep \*.json); do jsonlint \$i; done

Cron

show all crons for a user

crontab -l -u <username>

edit crons for your user

crontab -e

execute cron manually

run-parts /var/spool/cron

Test Email

yum install mailx

mail -s "test email" user@company.com < /dev/null

Centos Xauthority (graphical gui)

yum install -y xorg-x11-server-Xorg xorg-x11-xauth xorg-x11-apps

grep -i X11Forwarding /etc/ssh/sshd\_config (should be set to Yes)

ssh to box

ssh -X name@box

xclock (test)

SYSCTL

show all current values

sysctl -a

write new value

sysctl -w vm.swappiness=2

load values from file

sysctl -p /etc/sysctl.conf

monitor a command (run command repeatedly)

watch -n 5 free -h (runs free -h every 5 sec)

ENCRYPTION

OpenSSL

check cert expiration with OpenSSL

cat mycert.crt | openssl x509 -noout -enddate

Encrypt a file

openssl enc -aes-256-cbc -salt -in myFileUnencrypted.txt -out myFileEncrypted.txt.enc -k myPASSWORD

Unencrypt File

openssl enc -aes-256-cbc -in MyFileEncrypted.txt.enc -out myFileUnencrypted.txt <type in password>

7ZIP

install 7zip on centos

wget

https://www.mirrorservice.org/sites/dl.fedoraproject.org/pub/epel/7/x86\_64/Packag

es/p/p7zip-16.02-10.el7.x86\_64.rpm

wget

https://www.mirrorservice.org/sites/dl.fedoraproject.org/pub/epel/7/x86\_64/Packag

es/p/p7zip-plugins-16.02-10.el7.x86\_64.rpm

sudo rpm -U --quiet p7zip-16.02-10.el7.x86\_64.rpm

sudo rpm -U --quiet p7zip-plugins-16.02-10.el7.x86\_64.rpm

Encrypt a file

7za a -tzip -p -mem=AES256 testfile.zip testfile (enter password)

Decrypt a file

7za e testfile.zip (enter password)

GPG

send an encrypted file to a recipient

1. generate new gpg key pair

sender> gpg --gen-key

enter information including password and email address, password=S3nD3R

fred.sender@sender.com

if process is hanging on entropy, run the following to speed it up,

haveged -n 50g -f - | dd of=/dev/null

if password part is failing (may be due to TTY bug), run like this

script /dev/null

gpg --gen-key

2. Export the pub key to the recipient

sender> gpg --armor --output mypubkey.gpg --export fred.sender@sender.com

email this pub key to recipient

3.The reciever needs to generate their own public key w their own password and email

Import the recipient's pubkey into your GPG chain (have the recipient send you their pub key and provide the email associated to the pub key)

sender> gpg --import recipient-pubkey.gpg

have the recipient import sender's pub key the same way

4. Encrypt the file (has to be 1 single file, not multiple files or directories)

encrypt using senders private key

sender> gpg --output myFile.txt.gpg --encrypt --recipient joe.recipient@recipient.com myFile.txt

this will generate a binary gpg file

5. create a file signature (Checksum verification)

sender> shasum -a 256 myFile.txt | awk '{print \$1}' > myFile.txt.sha256sum

sender> gpg --output myFile.txt.sha256sum.sig --sign myFile.txt.sha256sum (enter password: S3nD3R)

(if shasum isnt installed, install with yum install -y perl-Digest-SHA)

email both binary gpg and .sig files to the recipient

6. Receiver unlocks the GPG file using the Reciever's password

receiver> gpg --output myFile.txt --decrypt myFile.txt.gpg (enter receiver password)

7. Verify the signature

receiver> gpg --output myFile.txt.sha256sum --decrypt myFile.txt.sha256sum.sig

additional

check pub GPG keys on host

gpg --list-keys --keyid-format LONG --fingerprint

check priv keys

gpg --list-secret-keys

delete pub key from keyring

gpg --delete-key D7B5FB7A (should be something like 2048R/D7BF5B7A)

delete priv key from keyring

gpg --delete-secret-key "Key name"

check which GPG key was used to encrypt a file

gpg --list-packets file.gpg

Resize a logical partition

Expand partition

1. add space to Hard Disk on VM in vCenter or VirtualBox

2. check all partitions, need to resize /opt its 30% full,

[root@mrxsplunkidx02 joe]# df -h

Filesystem Size Used Avail Use% Mounted on

/dev/mapper/vg0-root 7.8G 1.1G 6.3G 15% /

devtmpfs 1.9G 0 1.9G 0% /dev

tmpfs 1.9G 0 1.9G 0% /dev/shm

tmpfs 1.9G 8.6M 1.9G 1% /run

tmpfs 1.9G 0 1.9G 0% /sys/fs/cgroup

/dev/sda1 976M 110M 799M 13% /boot

/dev/mapper/vg0-home 2.0G 7.0M 1.8G 1% /home

This site uses cookies from Google to deliver its services and to analyze traffic. Information about your use of this site is shared with Google. By using this site, you agree to its use of cookies.

LEARN MORE GOT IT





mrxpalmeiras

- Home
- Wireshark cheat sheet
- Kail Pentest Netsec
- Rust
- Ansible
- Crystal
- Docker
- Elastic
- KDB Cheatsheet
- Git
- Golang
- JavaScript
- Linux
  - Iptables Basics
  - Linux Cheat Sheet
  - OpenVPN - Client config
  - Hi Frequency/Volume Trading - OS Tuning
  - SSH Certificate-based Authentication
  - Troubleshooting frozen system
  - RAID levels
  - Barrier - screen control across physical devices
  - QEMU Virtualization
  - TCP troubleshooting
  - Marten web framework
  - Graylog
  - htmx & hyperscript
  - UDP packet loss troubleshooting
- SSD types and terminology
- Data Structures
- More...
  - Julia cheatsheet
  - Sonicwall
  - Ninja Tools
  - Python
  - Ruby
  - Ruby Cheat Sheet
  - SaltStack
  - Scratchpad
  - Vagrant
  - Vue JS Cheatsheet

onfs sectors 0: 1 * 512 = 512 bytes					
Sector size (logical/physical): 512 bytes / 512 bytes					
I/O size (minimum/optimal): 512 bytes / 512 bytes					
Disk label type: dos					
Disk identifier: 0x000591a7					
	Device	Boot	Start	End	Blocks
	/dev/sda1	*	2048	1026047	512000
	/dev/sda2		1026048	22042623	10508288
					Id System
					83 Linux
					8e Linux LVM

Get ID and type of disk  
blkid

Additional Disk checks  
http://www.foxhop.net/local-or-san-device-in-linux

check what type of file system is on a partition  
lsblk -f

check disk for errors  
fsck (only works for certain filesystems)

check what kind of filesystem type  
df -T

show volume groups  
vgdisplay

extend volume group  
lvextend -r -L+25GB /dev/ivol/name

Mount a NetApp device as a local filesystem  
mount -t nfs -o \_netdev,rw,hard,intr,nosuid,dev,bg,nfsvers=3 netappNas01:/netbackup /netbackup

add to /etc/fstab,  
netappNas01:/netbackup /netbackup nfs  
\_netdev,rw,hard,intr,nosuid,dev,bg,nfsvers=3 0 0

check if Disks are local or mounted SAN  
ls /dev/disk/by-path/ (SANs will have an IP next to path)

Increase partition space via vCenter GUI  
Problem: current /opt only has 75G of available space, need to add another 20G  
df -h  
/dev/mapper/vg0-opt 80G 1.9G 75G 3% /opt

add disk space in vCenter console, increasing disk from 100GB to 120GB

on Centos box check name of scsi device,  
ls /sys/class/scsi\_device/  
0:0:0:0

rescan scsci bus  
echo 1 > /sys/class/scsi\_device/0\:0\:0\:0/device/rescan

check to see if extra space is visible,  
fdisk -l  
Disk /dev/sda: 128.8 GB

fdisk /dev/sda  
type 'p' - prints out all partitions  
type 'n' - create new partition  
type 'p' - to make new partition

select the next available sector (default), select default Last Sector

type 'w' to save changes  
reboot the VM  
once rebooted, type 'fdisk -l', a new partition is added

/dev/sda1	*	2048	2099199	1048576	83	Linux
/dev/sda2		2099200	2508799	204800	6	FAT16
/dev/sda3		2508800	209715199	103603200	8e	Linux LVM
/dev/sda4		209715200	251658239	20971520	83	Linux

now extend your /dev/mapper/vg0-opt

```
> vgs
VG #PV #LV #SN Attr VSize VFree
vg0 1 6 0 wz--n- 98.78g 0

> vgextend vg0 /dev/sda4
Volume group "vg0" successfully extended
```

Check to see available PE space (shows 20G of available space)  
> vgdisplay  
Free PE / Size 639 / <19.97 GiB

now resize to full available space, will show 94G of available space

```
> lvextend -l +100%FREE /dev/mapper/vg0-opt
> resize2fs /dev/mapper/vg0-opt
> df -h
/dev/mapper/vg0-opt 100G 1.9G 94G 2% /opt
```

### MDADM - Software RAID

Create RAID1 with btrfs on 2 physical disks on Centos 7  
check disks

lsblk					
	NAME	MAJ:MIN	RM	SIZE	RO TYPE MOUNTPOINT
	sda	8:0	0	40G	0 disk
	└─sda1	8:1	0	40G	0 part /
	sdb	8:16	0	8G	0 disk
	sdс	8:32	0	8G	0 disk

remove any existing partitions  
dd if=/dev/zero of=/dev/sdb bs=512 count=1

partition each disk (if using entire disk, can skip this entire Partition section)

```
fdisk /dev/sdb (do same with /dev/sdc)
n (new partition)
p (primary), select 1, enter, enter
t (select for RAID type), enter "fd"
w (write)
```

RAID examine:  
mdadm --examine /dev/sd[b-c]

Create RAID1  
mdadm --create /dev/md1 --level=mirror --raid-devices=2 /dev/sd[b-c]1  
(if no partitions present, remove 1 at end)

```
mdadm: Note: this array has metadata at the start and
may not be suitable as a boot device. If you plan to
store '/boot' on this device please ensure that
your boot-loader understands md/v1.x metadata, or use
--metadata=0.90
Continue creating array? y
mdadm: Defaulting to version 1.2 metadata
mdadm: array /dev/md1 started.
```

```
check RAID

mdadm --detail /dev/md1
/dev/md1:

    Version : 1.2
  Creation Time : Fri Aug 28 17:23:24 2020
    Raid Level : raid1
    Array Size : 8382464 (7.99 GiB 8.58 GB)
  Used Dev Size : 8382464 (7.99 GiB 8.58 GB)
    Raid Devices : 2
   Total Devices : 2
 Persistence : Superblock is persistent

               Update Time : Fri Aug 28 17:24:06 2020
                 State : clean
   Active Devices : 2
 Working Devices : 2
  Failed Devices : 0
   Spare Devices : 0
```

/dev/mapper/vg0-opt 4.8G 1.4G 3.3G 30% /opt

```
3. check logical space
lvs
appl vg0 -wi-ao---- 10.00g
home vg0 -wi-ao---- 2.00g
opt vg0 -wi-ao---- 5.00g
root vg0 -wi-ao---- 8.00g
swap vg0 -wi-ao---- 2.00g

check available HD space,
vgdisplay
Free PE / Size 191 / 5.97 GiB

4. Need to add another 5 Gigs to /opt
lvextend -r -L +5G /dev/mapper/vg0-opt

to extend ALL remaining free space,
lvextend -l +100%FREE /dev/mapper/vg0-opt

5. check the File System type of /opt
mount | grep opt
/dev/mapper/vg0-opt on /opt type ext4 (rw,relatime,data=ordered)

6. extend physical space
resize2fs /dev/mapper/vg0-opt

7. check space again, its now 15% full
df -h
/dev/mapper/vg0-opt 9.8G 1.4G 8.0G 15% /opt

check logical volume again,
lvs
opt vg0 -wi-ao---- 10.00g
```

Shrink Partition  
need to shrink partition /appl from 2GB to 1GB

```
lvs

appl vg0 -wi-ao---- 2.00g
docker vg0 -wi-ao---- 10.00g
home vg0 -wi-ao---- 60.00g
opt vg0 -wi-ao---- 5.00g
root vg0 -wi-ao---- 8.00g
swap vg0 -wi-ao---- 2.00g
tmp vg0 -wi-ao---- 3.00g
var vg0 -wi-ao---- 3.00g
```

unmount it  
umount -v /appl

get filesystem name  
df -h  
/dev/mapper/vg0-appl /appl

check for file system error  
e2fsck -ff /dev/mapper/vg0-appl  
(must pass all 5 stages)

reduce FS by 1GB  
resize2fs /dev/mapper/vg0-appl 1G

reduce the logical volume  
lvreduce -L -1G /dev/mapper/vg0-appl

mount /appl back on  
mount /dev/mapper/vg0-appl /appl

LVM  
check size of partition  
lvdisplay /appl

```
--- Logical volume ---
LV Path                /dev/vg0/appl
LV Name                appl
VG Name                vg0
LV UUID                Aim8Q2-gxp2-jnT0-0cS2-d3To-n5Nd-IJmvxo
LV Write Access         read/write
LV Creation host, time xxxx, 2018-02-23 11:52:48 -0500
LV Status                available # open 1
LV Size                1.00 GiB Current LE 32 Segments 1
Allocation              inherit
Read ahead sectors      auto - currently set to 8192
Block device            253:6
```

Remove Swap LV and merge it into Root LV  
want to remove 4GB swap LV and merge it into root, to give root more space,  
/dev/mapper/centos-root 50G 1.2G 49G 3% /

1. unmount and deactivate Swap LV  
lvchange -a n /dev/mapper/centos-swap

2. remove it  
lvremove /dev/mapper/centos-swap

3. extend Root volume  
lvextend -l +100%FREE /dev/mapper/centos-root

4. Grow the Root volume  
resize2fs /dev/mapper/centos-root  
(if XFS filesystem, use xfs\_growfs /dev/mapper/centos-root)

rename logical volume group  
vgdisplay (show all groups)

rename group centos to 'hc'  
vgrename /dev/centos /dev/hc

rename logical volume  
lvrename /dev/hc/disk1 /dev/hc/disk2

Remove Swapfile from /home and create new Swap LVM  
swapoff -a

Mount an EC2 volume as /home

attach volume to instance  
on ec2:  
lsblk

should be listed as nvme1n1 or similar name

check if filesystem has data  
file -s /dev/nvme1n1

if shows 'data', means volume is empty

create new volume  
mkfs -t xfs /dev/nvme1n1

create new mount point  
mkdir /home2

mount filesystem  
mount /dev/nvme1n1 /home2  
mv /home to /home\_old, move /home2 to /home

permanent mount, get ID of volume  
blkid (get UUID)

vi /etc/fstab  
UUID=<insert ID> /home xfs defaults 0 0

remount  
mount -a

Mount volume on EC2 as swap  
create volume, gp3 max iops (swap has to have fast read/write speed)  
attach volume to instance  
on instance,

```
lsblk (get new device name)
fdisk /dev/nvme3n1
n # new partition
t # partition type
82 # swap hex code
```

This site uses cookies from Google to deliver its services and to analyze traffic. Information about your use of this site is shared with Google. By using this site, you agree to its use of cookies.

LEARN MORE GOT IT





mrxpalmearas

- Home
- Wireshark cheat sheet
- Kali Pentest Netsec
- Rust
- Ansible
  - Crystal
- Docker
  - Elastic
  - KDB Cheatsheet
- Git
- Golang
- JavaScript
- Linux
  - Iptables Basics
  - Linux Cheat Sheet
  - OpenVPN - Client config
  - Hi Frequency/Volume Trading - OS Tuning
  - SSH Certificate-based Authentication
  - Troubleshooting frozen system
  - RAID levels
  - Barrier - screen control across physical devices
  - QEMU Virtualization
  - TCP troubleshooting
  - Marten web framework
  - Graylog
  - htmx & hyperscript
  - UDP packet loss troubleshooting
  - SSD types and terminology
  - Data Structures
- More...
  - Julia cheatsheet
  - Sonicwall
  - Ninja Tools
- Python
- Ruby
  - Ruby Cheat Sheet
- SaltStack
  - Scratchpad
- Vagrant
  - Vue JS Cheatsheet

```
Consistency Policy : resync

Name : min1:1 (local to host min1)
UUID : fae3d35a:862d0521:eb39400c:b9a794f0
Events : 17

Number  Major  Minor  RaidDevice State
0        8      17      0      active sync  /dev/sdb1
1        8      33      1      active sync  /dev/sdc1
```

create BTRFS filesystem

```
mkfs.btrfs /dev/md1
```

mount the filesystem

```
mount /dev/md1 /home (or mountpoint)
```

add to /etc/fstab

```
/dev/md1 /home btrfs defaults 0 0
```

create RAID config file

```
mdadm --detail --scan -v > /etc/mdadm.conf
```

test RAID1 by simulating drive failure

```
mdadm --manage --set-faulty /dev/md1 /dev/sdc1
```

check RAID status (will show DEGRADED)

```
mdadm --detail /dev/md1
```

setup alerts for Disk failure, add to /etc/mdadm.conf

```
MAILTO <your_email_addr>
DEVICE partitions
```

put scan in daemon

```
mdadm --monitor --scan --daemonize
```

add to kernel to start on boot

```
vi /etc/rc.local
```

add to bottom

```
/sbin/mdadm --monitor --scan --daemonize
```

add disk to array as hotspare

```
mdadm --grow /dev/md1 --add /dev/sdg1
```

add disk as full device

```
mdadm --grow /dev/md1 --add /dev/sdg1 --raid-devices=3
```

remove disk from array

```
mdadm /dev/md1 --fail /dev/sdc1 --remove /dev/sdc1
mdadm --grow /dev/md1 --raid-devices=2
```

use full size of array

```
mdadm --grow /dev/md1 --size=max
resize2fs /dev/md1
```

stop RAID

```
mdadm --stop /dev/md0
```

start RAID

```
mdadm --assemble /dev/md0 /dev/sdb1 /dev/sdc1
```

```
w # save
```

```
mkswap /dev/nvme3n1p1 # partition name
```

get the UUID of swap

blkid (find the partition UUID)

add to /etc/fstab

```
UUID=<UUID number> swap swap default 0 0
```

```
swapon /dev/nvme3n1p1
mount -a
```

BTRFS

check file on inode

```
btrfs inspect-internal inode-resolve 154326924 /mnt/hc
```

Extend disk space on EC2 instance (T3)

'/' currently at 94% usage

```
/dev/xvda1

go to EC2 console, click on instance > attached volumes > click on Volume >
modify > expand, add the additional disk space

root@host> lsblk
xvda TYPE=disk
---xvda1 TYPE=part /

growpart /dev/xvda 1

yum install xfsprogs

xfs_growfs -d / ## will resize / partition to full
```

Check disk health for bad sectors

- get disk name

```
lsblk | grep disk
```
- check blocks (read only)

```
sudo badblocks -v /dev/sda > badsector.txt
```

check disk read/write speed

```
hdparm -tv --direct /dev/sda
```

This site uses cookies from Google to deliver its services and to analyze traffic. Information about your use of this site is shared with Google. By using this site, you agree to its use of cookies.

LEARN MORE GOT IT



mrxpalmeiras

- Home
- Wireshark cheat sheet
- Kail Pentest Netsec
- Rust
- ▼ Ansible
- Crystal
- ▼ Docker
- Elastic
- KDB Cheatsheet
- ▼ Git
- ▼ Golang
- JavaScript
- ^ Linux
  - Iptables Basics
  - Linux Cheat Sheet**
  - OpenVPN - Client config
  - Hi Frequency/Volume Trading - OS Tuning
  - SSH Certificate-based Authentication
  - Troubleshooting frozen system
  - RAID levels
  - Barrier - screen control across physical devices
  - QEMU Virtualization
  - TCP troubleshooting
  - Marten web framework
  - Graylog
  - htmx & hyperscript
  - UDP packet loss troubleshooting
- SSD types and terminology
- Data Structures
- ▼ More...
- Julia cheatsheet
- Sonicwall
- Ninja Tools
- ▼ Python
- ▼ Ruby
- Ruby Cheat Sheet
- ▼ SaltStack
- Scratchpad
- ▼ Vagrant
- Vue JS Cheatsheet

Regex / awk / sed / tr

find all lines starting with #  
^#.\*\$  
  
find blank line  
^\s\*\$  
  
  
remove leading and trailing commas from string  
str=",there was a loud, bang, there,"  
str="\${str#},""  
str="\${str%},""  
echo \$str >> "there was a loud, bang, there"

SED

Replace string in a file (write directly to file -i)  
sed -i -e "s/\${prev\_version}/\${version}/g" bitbucket.service  
  
Replace anything between 2 delimiters "!!" with word "super"  
sed -e 's/!.\*!/super/g' /etc/file  
  
remove whitespace  
sed -i "s/ //g" file # replace inline  
sed "s/[[:space:]]//g" file # replace just on screen  
  
remove 2nd line from top, from file  
sed -i '2,\$d' file  
  
replace newline with comma  
sed ':a;N;\$!ba;s/\n/ /g'  
  
Remove leading spaces and tabs  
sed 's/^[ \t]\*//'  
  
Remove single spaces only (leave multiple spaces)  
sed 's/\(.\) //g'  
  
Reduce multiple spaces to one  
sed 's/ \+/ /g'  
  
Replace multiple newlines with a single newline  
sed '/^\$/N;/^\\n\$/D' file.txt  
  
Delete text in a line between two markers {}  
sed -e 's/(.{\\}).\*\\(\\)///'  
  
Remove empty lines  
sed '/^\\s\*\$d'  
  
Remove all but the first line matching pattern  
sed '2,\${/pattern/d;}'  
  
Remove only the first line matching pattern  
sed '0,/pattern/{/pattern/d;}'  
  
Move the first line to the end of the list  
sed '1,1{H;1h;d;};\$G'  
  
Remove non-alphanumeric characters from words  
sed 's/^[[:alnum:]]-\\ //g'  
  
Reduce multiple spaces to one for a line containing a string  
iostat | sed -n '/^sd/s/ \+/ /gp'  
  
insert "apple" into beginning of a file  
sed 's/^/apple /' file1  
  
insert apple into end of file  
sed 's/\$/apple /' file1  
  
replace orange with apple, only on 3rd line of text  
sed '3s/orange/apple/g' file1  
  
replace orange with apple from 1st to 3rd line  
sed '1,3s/orange/apple/g' file1  
  
replace multiple words, orange with apple, red with blue  
sed 's/orange/apple/g; s/red/blue/g' file1  
  
remove 1st occurrence of specific character "b"  
sed 's/b//' file1  
  
remove all instances of "b" in file  
sed 's/b//g' file1  
  
remove last character of every line  
sed 's/.\$//' file1  
  
remove "b" only if its last character in line  
sed 's/b\$//' file1  
  
remove all numbers in every line of a file  
sed 's/[0-9]//g' file1

TR

remove whitespace from string  
fields=\$(echo -e "\${fields}" | tr -d '[:space:]')  
  
or use xargs  
echo \$fields | xargs (will strip leading and trailing whtiespace)  
  
strip double quotes  
echo '"string"' | tr -d '"'  
  
Identify server's primary IP address  
/sbin/ifconfig | sed -m 's/127.0.0.1//s/\\.inet (addr:)?([0-9]\\.[0-9]\*).\*/p'  
  
split a string by a delimiter  
string="apple, cherry, banana"  
first=\$(echo \$string | cut -d',' -f1)  
second=\$(echo \$string | cut -d',' -f2)  
  
Remove non-printable characters from files  
tr -cd '6' < infile > outfile

XARGS

find all files with JPG extension and rename each file to be JPEG2 (xargs -i)  
ls | grep ".jpg" | xargs -I {} mv {} {}.jpeg2

examples taken from: <https://www.igroseyedko.com/awk-sed-snippets-for-sysadmins/>

AWK

Built-In AWK functions [full list of functions](#)

to Uppercase  
awk -F: '{print toupper(\$1)}' file.txt  
NAME: JOE  
  
print entire line (\$0)  
joe:employees:123  
bob:employees:222  
awk -F":" '{print \$0}' file.txt  
  
Awk If-Else  
awk -F: '{if(\$1=="name") print \$2;else print "NONE"}' file.txt  
joe  
  
find by Regex  
awk -F: '/ing\$/ ' file.txt  
status: thinking  
occupation: moving

```
print all values that match 1st column = "color"
cat file.txt
color:red
size:25
color:blue
size:50

awk -F: '$1 ~ /color$/ {print $2}' file.txt
red
blue
```

split string by delimiter  
third=\$(echo \$string | awk -F",", '{print \$3}')

or use delimiter flag  
echo \$string | awk '{print \$1,\$3}' FS=","  
  
Remove commas inside double-quotes  
awk -F'"' -v OFS=" '{ for (i=2; i<=NF; i+=2) gsub(",","",\$i)} 1'

Remove duplicate words in a line  
awk '{ while(++i<=NF) printf (!a[i]++) ? \$i FS : " "; i=split("",a); print "" }'

Remove duplicate lines in a file without sorting  
cat file | awk '!a[\$0]++'

Print number of characters for each line in a file  
awk '{ print length(\$0)"\t"\$0; }' file.txt

Begin and End Functions  
awk -F: 'BEGIN {print "this is beginning"} {print \$0} END {print "End!"}' file.txt  
this is beginning  
name: joe  
End!

remove all duplicate entries from a file  
awk '!x[\$0]++' filename

```
cat file.txt
apple,300
grape,200
grape,400
apple,500
banana,200

print only lines with word "apple"
awk '/apple/' file.txt

print only "grape" record, print 2nd column only
awk '$0 ~ /grape/{print $2}' file.txt

print any line that does not contain "apple"
awk '!/apple/' file.txt

print any line that has grape or banana
awk -F, '$1 ~ /^grape|^banana/' file.txt

print any line where number is greater than 200
awk -F, '$2>200' file.txt

print any line thats greater than 200 or has "grape"
```

Remove entire words containing non-alphabetic characters  
awk '{ofs=""; for (i=1; i<=NF; i++) if (\$i ~ /^[[:alpha:]]+\$/){printf "%s%s", ofs, \$i; ofs=OFS} print "" }'

-----  
  
Sample "temp" file  
ID1,223  
ID2,124  
ID3,125  
ID2,400  
  
Add up values in second column  
awk -F",", '{s+=\$2}END{print s}' temp  
  
Add up the values in the second column only for ID2  
awk -F, '\$1=="ID2"{s+=\$2;}END{print s}' temp  
v="ID2"; awk -F, -v v="\$v" '\$1==v{s+=\$2;}END{print s}' temp  
  
List unique values in 1st column  
awk -F, '{a[\$1];}END{for (i in a)print i;}' temp  
  
Add up values in the second column for each ID  
awk -F, '{a[\$1]+=\$2;}END{for(i in a)print i, "a[i];}' temp  
  
Remove only the first line matching pattern  
awk '!/pattern/ || f++'  
  
Remove all but the first line matching pattern  
awk '/pattern/&&f++ {next} 1'  
  
Show allocated disk space  
df -k1P -t xfs -t ext2 -t ext3 -t ext4 -t reiserfs | grep -oE '[0-9]{1,}([+][0-9]{1,})+' | awk '{sum\_used += \$1} END {printf "%.0f GB\n", sum\_used/1024/1024}'

This site uses cookies from Google to deliver its services and to analyze traffic. Information about your use of this site is shared with Google. By using this site, you agree to its use of cookies.

LEARN MORE GOT IT



mrxpalmeiras

- Home
- Wireshark cheat sheet
- Kali Pentest Netsec
- Rust
- Ansible
- Crystal
- Docker
- Elastic
- KDB Cheatsheet
- Git
- Golang
- JavaScript
- Linux
  - Iptables Basics
  - Linux Cheat Sheet
  - OpenVPN - Client config
  - Hi Frequency/Volume Trading - OS Tuning
  - SSH Certificate-based Authentication
  - Troubleshooting frozen system
  - RAID levels
  - Barrier - screen control across physical devices
  - QEMU Virtualization
  - TCP troubleshooting
  - Marten web framework
  - Graylog
  - htmx & hyperscript
  - UDP packet loss troubleshooting
- SSD types and terminology
- Data Structures
- More...
  - Julia cheatsheet
  - Sonicwall
  - Ninja Tools
- Python
- Ruby
  - Ruby Cheat Sheet
- SaltStack
- Scratchpad
- Vagrant
- Vue JS Cheatsheet

## Sysdig

install sysdig  
curl -s https://s3.amazonaws.com/download.draios.com/stable/install-sysdig | sudo bash

write to scap file, 200MB in size, keep only 5 files  
sysdig -C 200 -W 5 -w dump.scap

search for specific port  
sysdig fd.port=8335

search for specific process  
sysdig proc.name=sudo

see every action done by every user  
sysdig -c spy\_users

### Sysdig Network

see top procs in network bandwidth usage  
sysdig -c topprocs\_net

show network data exchanged with host 192.168.38.5  
sysdig -s2000 -A -c echo\_fds fd.cip=192.168.38.5

show top client IPs in terms of established connections  
sysdig -c fdcount\_by fd.cip "evt.type=accept"

in terms of total bytes  
sysdig -c fdbytes\_by fd.cip

see all GET HTTP requests made  
sudo sysdig -s 2000 -A -c echo\_fds fd.port=80 and evt.buffer contains GET

### Sysdig disk IO

see top procs in terms of bandwidth usage  
sysdig -c topprocs\_file

list procs that use high number of files  
sysdig -c fdcount\_by proc.name "fd.type=file"

see top files in r/w bytes  
sysdig -c topfiles\_bytes

see top directories in terms of r/w disk activity  
sysdig -c fdbytes\_by fd.directory "fd.type=file"

in specific directory  
sysdig -c fdbytes\_by fd.filename "fd.directory=/tmp/"

observe IO activity on all files named "passwd"  
sysdig -A -c echo\_fds "fd.filename=passwd"

show top procs in terms of IO errors  
sysdig -c topprocs\_errors

Sysdig Proc and CPU

show top procs by CPU usage  
sysdig -c topprocs\_cpu

observe standard output of proc  
sysdig -s4096 -A -c stdout proc.name=myproc

Sysdig security

show all file opens in /etc directory  
sysdig evt.type=open and fd.name contains /etc

Show the ID of all the login shells that have launched the "tar" command  
sysdig -r file.scap -c list\_login\_shells tar

Show all the commands executed by the login shell with the given ID  
sysdig -r trace.scap.gz -c spy\_users proc.loginshellid=5459

## Snap

snap alias  
snap install python38  
snap alias python38 python3.8 (creates symlink /usr/bin/python3.8)

## Troubleshoot

### slow / frozen system

check if procs are in uninterrupted sleep state (waiting for IO and causing slowness)  
ps aux (check STAT column, will show procs that are in uninterrupted sleep)

check paging faults

sar -B 2 5 will generate paging report, check majflt column, major faults per second, if high #, means system is out of RAM

### SysRQ

enable SysRQ to kill procs that are in 'uninterrupted sleep' state. SysRQ will respond even in frozen state (assuming command line is responsive)

configure server to have sysrq enabled

echo 1 > /proc/sys/kernel/sysrq

add to sysctl

sysctl -w kernel.sysrq=1

- enable sysrq
- kill D state procs

Check why a server rebooted  
sudo ausearch -i -m system\_boot,system\_shutdown | tail -4

show all reboots  
journalctl --list-boots

### VMstat - CPU,Memory, I/O usage

check swap and memory allocation  
vmstat -S M 5 (updates every 5 seconds)

**Proc**  
r: The number of runnable processes. These are processes that have been launched and are either running or are waiting for their next time-sliced burst of CPU cycles.  
b: The number of processes in uninterruptible sleep. The process isn't sleeping, it is performing a blocking system call, and it cannot be interrupted until it has completed its current action. Typically the process is a device driver waiting for some resource to come free. Any queued interrupts for that process are handled when the process resumes its usual activity.

**Memory**  
swpd: the amount of virtual memory used. In other words, how much memory has been swapped out.,  
free: the amount of idle (currently unused) memory.  
buff: the amount of memory used as buffers.  
cache: the amount of memory used as cache.

**Swap**  
si: Amount of virtual memory swapped in from swap space.  
so: Amount of virtual memory swapped out to swap space.

**IO**  
bi: Blocks received from a block device. The number of data blocks used to swap virtual memory back into RAM.  
bo: Blocks sent to a block device. The number of data blocks used to swap virtual memory out of RAM and into swap space.

**System**  
in: The number of interrupts per second, including the clock.  
cs: The number of context switches per second. A context switch is when the kernel swaps from system mode processing into user mode processing.

**CPU**  
These values are all percentages of the total CPU time.  
us: Time spent running non-kernel code. That is, how much time is spent in user time processing and in nice time processing.  
sy: Time spent running kernel code.  
id: Time spent idle.  
wa: Time spent waiting for input or output.  
st: Time stolen from a virtual machine. This is the time a virtual machine has to wait for the hypervisor to finish servicing other virtual machines before it can come back and attend to this virtual machine.

## Augeas

example file content: /etc/pam.d/password-auth

auth	required	pam_env.so
auth	required	pam_faildelay.so delay=2000000
auth	[success=1 default=bad]	pam_unix.so nullok try_first_pass
auth	requisite	pam_succeed_if.so uid >= 1000 quiet_success
auth	required	pam_deny.so
account	required	pam_unix.so
account	sufficient	pam_localuser.so
account	sufficient	pam_succeed_if.so uid < 1000 quiet
account	required	pam_permit.so
password	requisite	pam_pwquality.so try_first_pass local_users_only retry=3
authtok_type=		
password	sufficient	pam_unix.so sha512 shadow nullok try_first_pass use_authtok
password	required	pam_deny.so
session	optional	pam_keyinit.so revoke
session	required	pam_limits.so
-session	optional	pam_systemd.so
session	[success=1 default=ignore]	pam_succeed_if.so service in crond quiet use_uid
session	required	pam_unix.so
<b>password sufficient pam_unix.so remember=15</b>		
auth	required	pam_faillock.so preauth audit silent deny=5 unlock_time=900
auth	[default=die]	pam_faillock.so authfail audit deny=5 unlock_time=900
auth	sufficient	pam_faillock.so authsucc audit deny=5 unlock_time=900
password	requisite	pam_pwquality.so try_first_pass retry=3

search through tree for specific string  
augtool> print /files/etc/pam.d/password-auth/\*[module='pam\_unix.so']/argument[. = 'remember=15']

search using a Regex  
augtool> print /files/etc/pam.d/password-auth/\*[module='pam\_unix.so']  
[type='password'][control='sufficient']/argument[. =~ regexp("remember=.\*")]

get last node of a tree  
augtool> print /files/etc/pam.d/password-auth/\*[module='pam\_unix.so']  
[type='password'][control='sufficient'][last()]

add Argument value on last node, last argument position  
set /files/etc/pam.d/password-auth/\*[module='pam\_unix.so'][type='password']  
[control='sufficient'][last()]/argument[last()+1] 'remember=33'

This site uses cookies from Google to deliver its services and to analyze traffic. Information about your use of this site is shared with Google. By using this site, you agree to its use of cookies.

LEARN MORE GOT IT





mrxpalmeiras

- Home
- Wireshark cheat sheet
- Kail Pentest Netsec
- Rust
- ▼ Ansible
- Crystal
- ▼ Docker
- Elastic
- KDB Cheatsheet
- ▼ Git
- ▼ Golang
- JavaScript
- ⋈ Linux
- Iptables Basics
- Linux Cheat Sheet
- OpenVPN - Client config
- Hi Frequency/Volume Trading - OS Tuning
- SSH Certificate-based Authentication
- Troubleshooting frozen system
- RAID levels
- Barrier - screen control across physical devices
- QEMU Virtualization
- TCP troubleshooting
- Marten web framework
- Graylog
- htmx & hyperscript
- UDP packet loss troubleshooting
- SSD types and terminology
- Data Structures
- ▼ More...
- Julia cheatsheet
- Sonicwall
- Ninja Tools
- ▼ Python
- ▼ Ruby
- Ruby Cheat Sheet
- ▼ SaltStack
- Scratchpad
- ▼ Vagrant
- Vue JS Cheatsheet

Q

/proc pseudofiles

/proc/PID/cmdline - Command line arguments.

/proc/PID/cpu - Current and last cpu in which it was executed.

/proc/PID/cwd - Link to the current working directory.

/proc/PID/environ - Values of environment variables.

/proc/PID/exe - Link to the executable of this process.

/proc/PID/fd - Directory, which contains all file descriptors.

/proc/PID/maps - Memory maps to executables and library files.

/proc/PID/mem - Memory held by this process.

/proc/PID/root - Link to the root directory of this process.

/proc/PID/stat - Process status.

/proc/PID/statm - Process memory status information.

/proc/PID/status - Process status in human readable form.

/proc/apm - Advanced power management info.

/proc/bus - Directory containing bus specific information.

/proc/cmdline - Kernel command line.

/proc/cpuinfo - Information about the processor, such as its type, make, model, and performance.

/proc/devices - List of device drivers configured into the currently running kernel (block and character).

/proc/dma - Shows which DMA channels are being used at the moment.

/proc/driver -Various drivers grouped here, currently rtc

/proc/execddomains - Execdomains, related to security.

/proc/fb - Frame Buffer devices.

/proc/filesystems - Filesystems configured/supported into/by the kernel.

/proc/fs - File system parameters, currently nfs/exports.

/proc/interrupts - Shows which interrupts are in use, and how many of each there have been.

/proc/iomem - Memory map.

/proc/loports - Which I/O ports are in use at the moment.

/proc/irq - Masks for irq to cpu affinity.

/proc/isapnp - ISA PnP (Plug&Play) Info.

/proc/kmsg - Messages output by the kernel. These are also routed to syslog.

/proc/ksyms - Kernel symbol table.

/proc/loadavg - The 'load average' of the system; three indicators of how much work the system has done during the last 1, 5 & 15 minutes.

/proc/locks - Kernel locks.

/proc/meminfo - Information about memory usage, both physical and swap. Concatenating this file produces similar results to using 'free' or the first few lines of 'top'.

/proc/misc - Miscellaneous pieces of information. This is for information that has no real place within the rest of the proc filesystem.

/proc/modules - Kernel modules currently loaded. Typically its output is the same as that given by the 'lsmod' command.

/proc/mounts - Mounted filesystems (same as running 'mount' cmd)

/proc/net - Status information about network protocols.

/proc/stat - Overall/various statistics about the system, such as the number of page faults since the system was booted.

/proc/swaps - Swap space utilization

/proc/sys - This is not only a source of information, it also allows you to change parameters within the kernel without the need for recompilation or even a system reboot. Take care when attempting this as it can both optimize your system and also crash it.

/proc/sys/fs - Contains file system data. This subdirectory contains specific file system, file handle, inode, dentry and quota information.

/proc/sys/vm - The files in this directory can be used to tune the operation of the virtual memory (VM) subsystem of the Linux kernel. In addition, one of the files (bdflush) has some influence on disk usage.

nfract - This parameter governs the maximum number of dirty buffers in the buffer cache. Dirty means that the contents of the buffer still have to be written to disk (as opposed to a clean buffer, which can just be forgotten about). Setting this to a higher value means that Linux can delay disk writes for a long time, but it also means that it will have to do a lot of I/O at once when memory becomes short. A lower value will spread out disk I/O more evenly.

ndirty - Ndirty gives the maximum number of dirty buffers that bdflush can write to the disk at one time. A high value will mean delayed, bursty I/O, while a small value can lead to memory shortage when bdflush isn't woken up often enough.

nrefill - This is the number of buffers that bdflush will add to the list of free buffers when refill\_freelist() is called. It is necessary to allocate free buffers beforehand, since the buffers are often different sizes than the memory pages and some bookkeeping needs to be done beforehand. The higher the number, the more memory will be wasted and the less often refill\_freelist() will need to run.

nref\_dirt - When refill\_freelist() comes across more than nref\_dirt dirty buffers, it will wake up bdflush.

age\_buffer, age\_super - Finally, the age\_buffer and age\_super parameters govern the maximum time Linux waits before writing out a dirty buffer to disk. The value is expressed in jiffies (clockticks), the number of jiffies per second is 100. Age\_buffer is the maximum age for data blocks, while age\_super is for filesystems meta data.

buffermem - The three values in this file control how much memory should be used for buffer memory. The percentage is calculated as a percentage of total system memory. The values are:

min\_percent - This is the minimum percentage of memory that should be spent on buffer memory.

borrow\_percent - When Linux is short on memory, and the buffer cache uses more than it has been allotted, the memory management (MM) subsystem will prune the buffer cache more heavily than other memory to compensate.

max\_percent - This is the maximum amount of memory that can be used for buffer memory.

freepages - This file contains three values: min, low and high:

min - When the number of free pages in the system reaches this number, only the kernel can allocate more memory.

low - If the number of free pages falls below this point, the kernel starts swapping aggressively.

high - The kernel tries to keep up to this amount of memory free; if memory falls below this point, the kernel starts gently swapping in the hopes that it never has to do really aggressive swapping.

kswapd - Kswapd is the kernel swap out daemon. That is, kswapd is that piece of the kernel that frees memory when it gets fragmented or full. Since every system is different, you'll probably want some control over this piece of the system. The file contains three numbers:

tries\_base - The maximum number of pages kswapd tries to free in one round is calculated from this number. Usually this number will be divided by 4 or 8 (see mm/vmscan.c), so it isn't as big as it looks. When you need to increase the bandwidth to/from swap, you'll want to increase this number.

tries\_min - This is the minimum number of times kswapd tries to free a page each time it is called. Basically it's just there to make sure that kswapd frees some pages even when it's being called with minimum priority.

swap\_cluster - This is probably the greatest influence on system performance. swap\_cluster is the number of pages kswapd writes in one turn. You'll want this value to be large so that kswapd does its I/O in large chunks and the disk doesn't have to seek as often, but you don't want it to be too large since that would flood the request queue.

overcommit\_memory - This file contains one value. The following algorithm is used to decide if there's enough memory: if the value of overcommit\_memory is positive, then there's always enough memory. This is a useful feature, since programs often malloc() huge amounts of memory 'just in case', while they only use a small part of it. Leaving this value at 0 will lead to the failure of such a huge malloc(), when in fact the system has enough memory for the program to run. On the other hand, enabling this feature can cause you to run out of memory and thrash the system to death, so large and/or important servers will want to set this value to 0.

pagecache - This file does exactly the same job as buffermem, only this file controls the amount of memory allowed for memory mapping and generic caching of files. You don't want the minimum level to be too low, otherwise your system might thrash when memory is tight or fragmentation is high.

pagetable\_cache - The kernel keeps a number of page tables in a per-processor cache (this helps a lot on SMP systems). The cache size for each processor will be between the low and the high value. On a low-memory, single CPU system, you can safely set these values to 0 so you don't waste memory. It is used on SMP systems so that the system can perform fast pagetable allocations without having to acquire the kernel memory lock. For large systems, the settings are probably fine. For normal systems they won't hurt a bit. For small systems ( less than 16MB ram) it might be advantageous to set both values to 0.

swaptctl - This file contains no less than 8 variables. All of these values are used by kswapd. The first four variables sc\_max\_page\_age, sc\_page\_advance, sc\_page\_decline and sc\_page\_initial\_age are used to keep track of Linux's page aging. Page ageing is a bookkeeping method to track which pages of memory are often used, and which pages can be swapped out without consequences.

/proc/sys/net/core - Network core options

rmem\_default - The default setting of the socket receive buffer in bytes.

rmem\_max - The maximum receive socket buffer size in bytes.

wmem\_default - The default setting (in bytes) of the socket send buffer.

wmem\_max - The maximum send socket buffer size in bytes.

message\_burst and message\_cost - These parameters are used to limit the warning messages written to the kernel log from the networking code. They enforce a rate limit to make a denial-of-service attack impossible. A higher message\_cost factor, results in fewer messages that will be written. Message\_burst controls when messages will be dropped. The default settings limit warning messages to one every five seconds.

netdev\_max\_backlog - Maximum number of packets, queued on the INPUT side, when the interface receives packets faster than kernel can process them.

optmem\_max - Maximum ancillary buffer size allowed per socket. Ancillary data is a sequence of struct cmsghdr structures with appended data.

This site uses cookies from Google to deliver its services and to analyze traffic. Information about your use of this site is shared with Google. By using this site, you agree to its use of cookies.

LEARN MORE GOT IT